

Introduction to Deep Learning with Keras

Monday, 17 August 2020 7:30 PM

Introducing keras

Theano vs Keras

```
import theano
import theano.tensor as T
from theano.ifelse import ifelse
import numpy as np
from random import random

# Define variables
x = T.matrix('x')
w1 = theano.shared(np.array([random(), random()]))
w2 = theano.shared(np.array([random(), random()]))
b1 = theano.shared(np.array([random(), random()]))
b2 = theano.shared(np.array([random(), random()]))

z1 = 1/(1+T.exp(-T.dot(x,w1)+b1))
x1 = T.tanh(z1*w1+1)
a1 = 1/(1+T.exp(-T.dot(x1,w2)+b2))

a2_hat = T.vector('a2_hat').astype(theano.config.floatX)
cost = - (a2_hat*T.log(a1) + (1-a2_hat)*T.log(1-a1)).mean()
dw1, dw2, db1, db2 = T.grad(cost, [w1, w2, b1, b2])

# You can [finally] train your model
cost = []
for i in range(10000):
    print cost
    print_theo = train(theo_inputs, theo_outputs)
    cost.append(cost_iter)
```

```
from keras.layers import Dense
from keras.models import Sequential

# Define model and add layers
model = Sequential()
model.add(Dense(2, input_shape=(2,), activation='sigmoid'))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='categorical_crossentropy')

# Train model
model.fit(inputs, outputs)
```

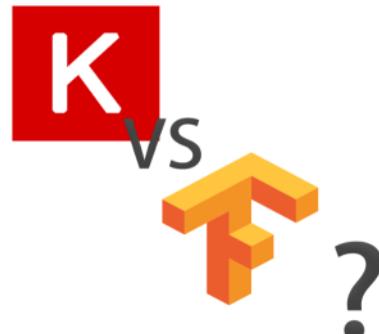
Keras

- Deep Learning Framework
- Enables fast experimentation
- Runs on top of other frameworks
- Written by François Chollet



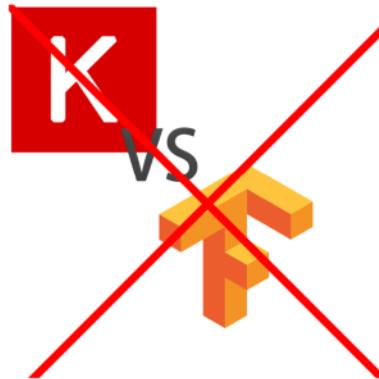
Why use Keras?

- Fast industry-ready models
- For beginners and experts
- Less code
- Build any architecture
- Deploy models in multiple platforms

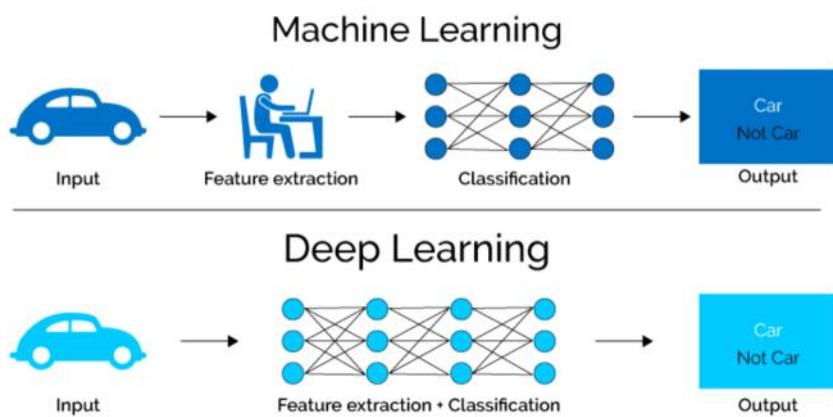


Keras + TensorFlow

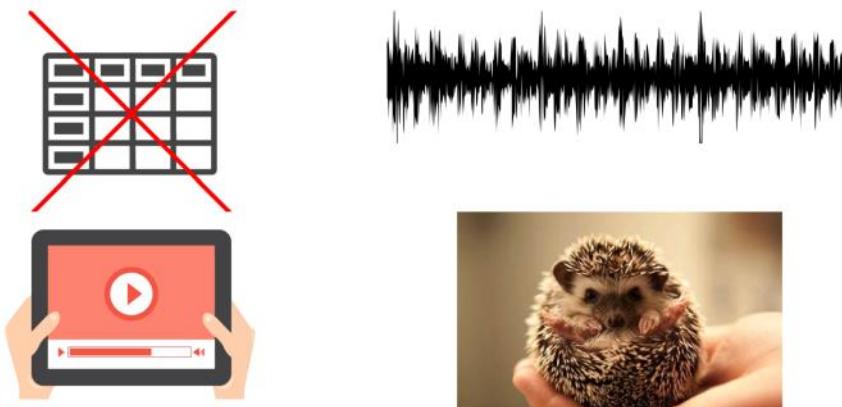
- TensorFlow's high level framework of choice
- Keras is complementary to TensorFlow
- You can use TensorFlow for low level features



Feature Engineering



Unstructured data

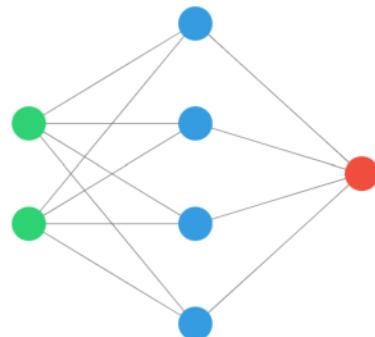


So, when to use neural networks?

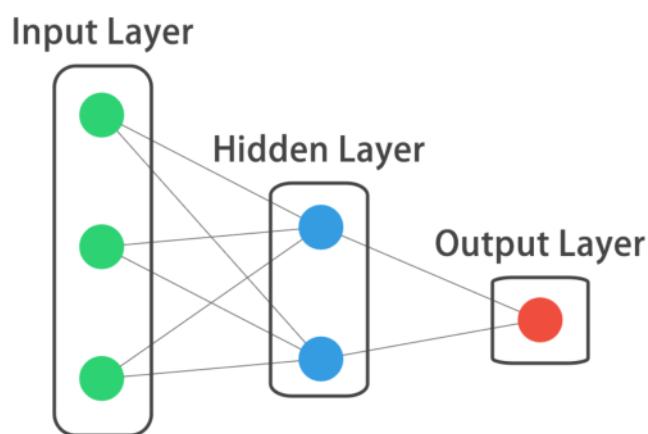
- Dealing with unstructured data
- Don't need easily interpretable results
- You can benefit from a known architecture

Example: Classify images of cats and dogs

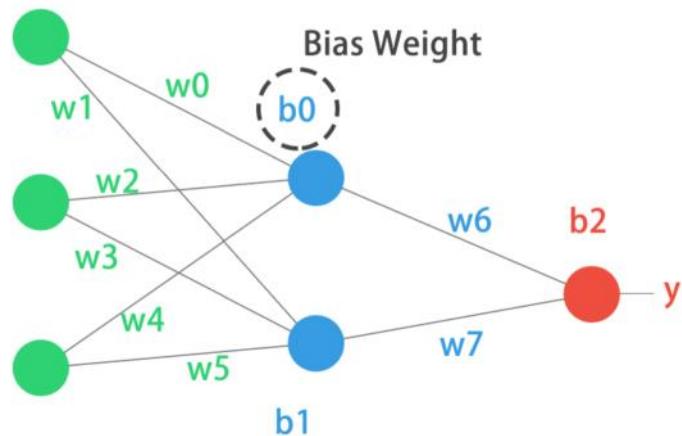
- Images \rightarrow Unstructured data
- You don't care about why the network knows it's a cat or a dog
- You can benefit from convolutional neural networks



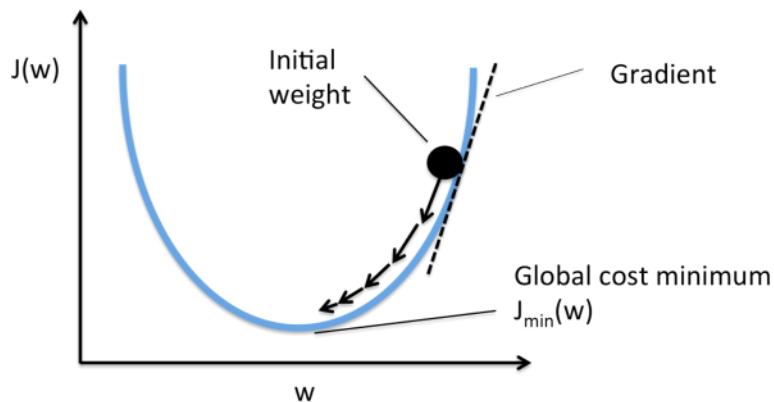
A neural network?



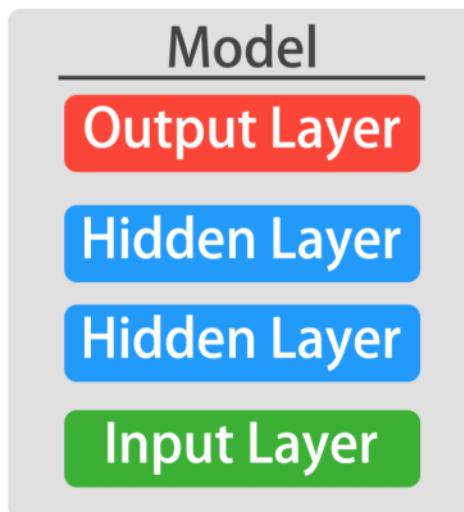
Parameters



Gradient descent



The sequential API



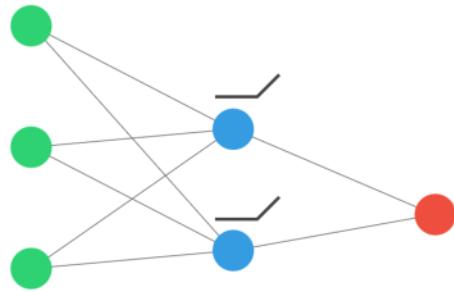
Adding activations

```
from keras.models import Sequential
from keras.layers import Dense

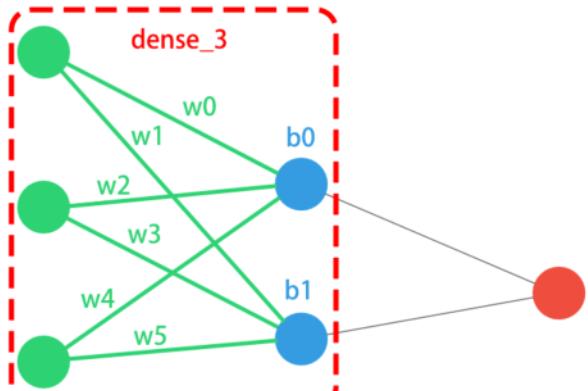
# Create a new sequential model
model = Sequential()

# Add an input and dense layer
model.add(Dense(2, input_shape=(3,),
               activation="relu"))

# Add a final 1 neuron layer
model.add(Dense(1))
```



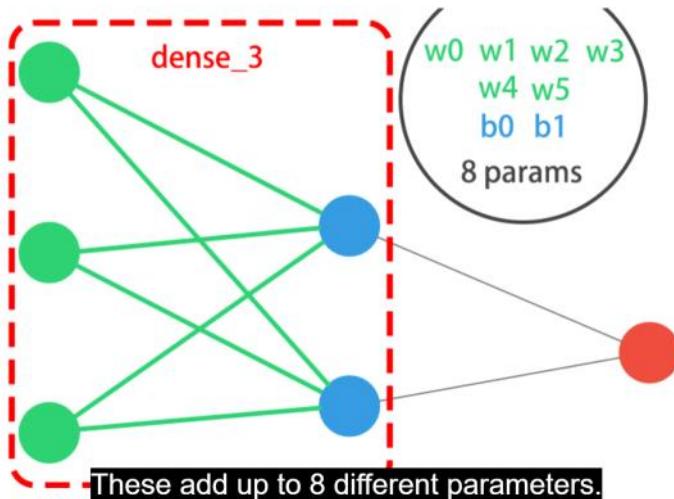
Visualize parameters



That's why we see that this layer has 8 parameters: 6 parameters or weights come

missing 2 parameters come from the bias weights, b0 and b1, 1 per each neuron in
the hidden layer.

Visualize parameters

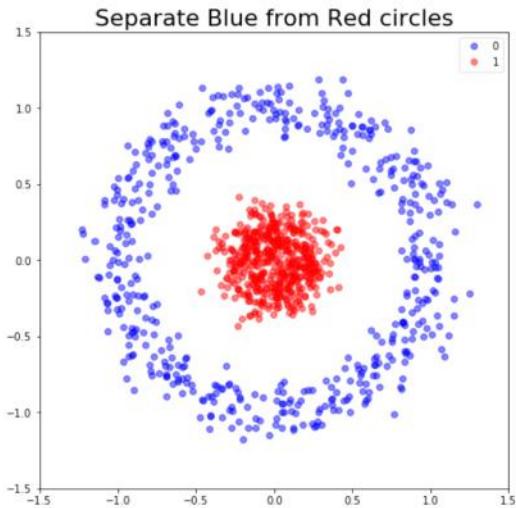


Summarize your model!

```
model.summary()
```

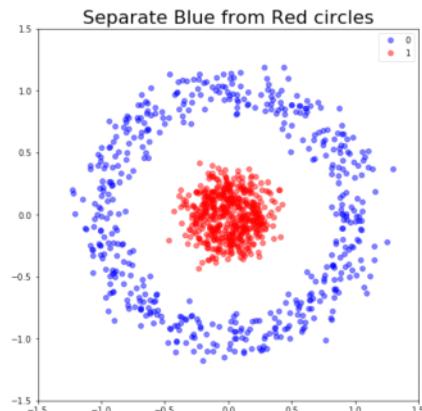
Layer (type)	Output Shape	Param #
<hr/>		
dense_3 (Dense)	(None, 2)	--> 8 <--
<hr/>		
dense_4 (Dense)	(None, 1)	3
<hr/>		
Total params:	11	
Trainable params:	11	
Non-trainable params:	0	

When to use binary classification?



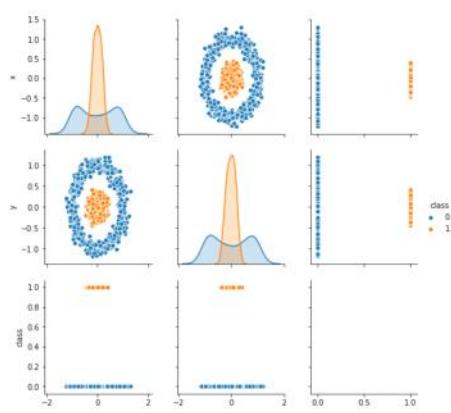
Our dataset

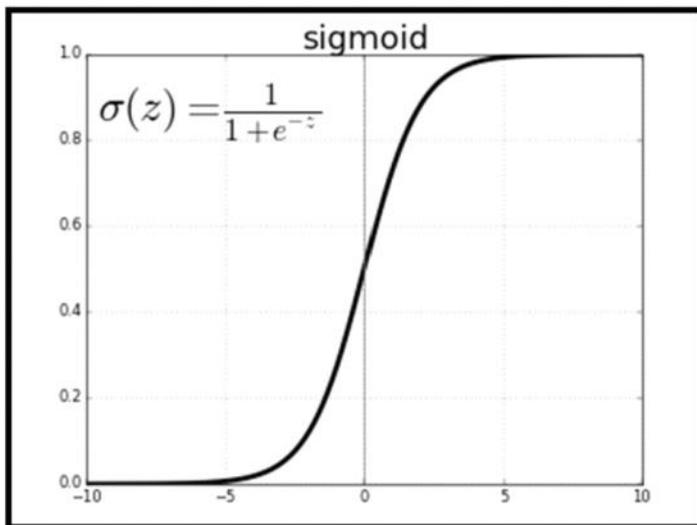
coordinates	labels
[0.242, 0.038]	1
[0.044, -0.057]	1
[-0.787, -0.076]	0



Pairplots

```
import seaborn as sns  
  
# Plot a pairplot  
sns.pairplot(circles, hue="target")
```





Output Layer

The sigmoid function

neuron output	$3 \rightarrow$	sigmoid	$\rightarrow 0.95 \rightarrow$	transformed output	rounded output
------------------	-----------------	---------	--------------------------------	-----------------------	-------------------

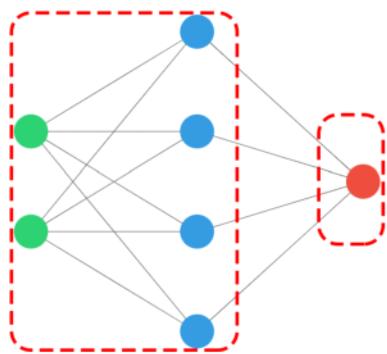
Let's build it

```
from keras.models import Sequential
from keras.layers import Dense

# Instantiate a sequential model
model = Sequential()

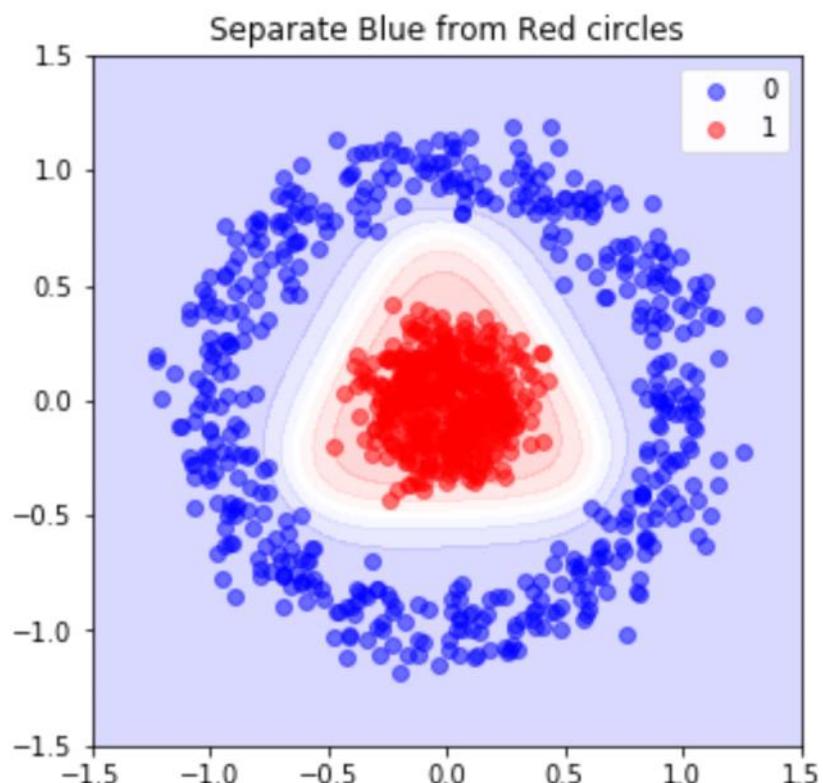
# Add input and hidden layer
model.add(Dense(4, input_shape=(2,), activation='tanh'))

# Add output layer, use sigmoid
model.add(Dense(1, activation='sigmoid'))
```

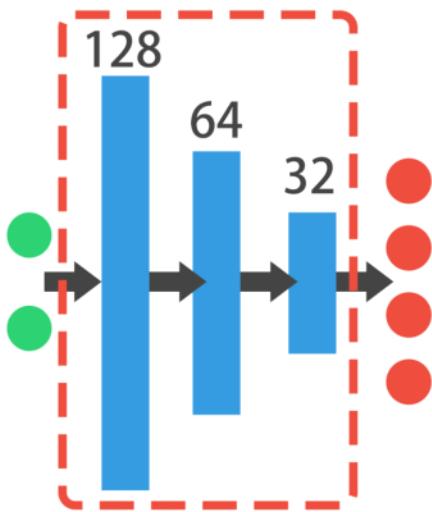


Compiling, training, predicting

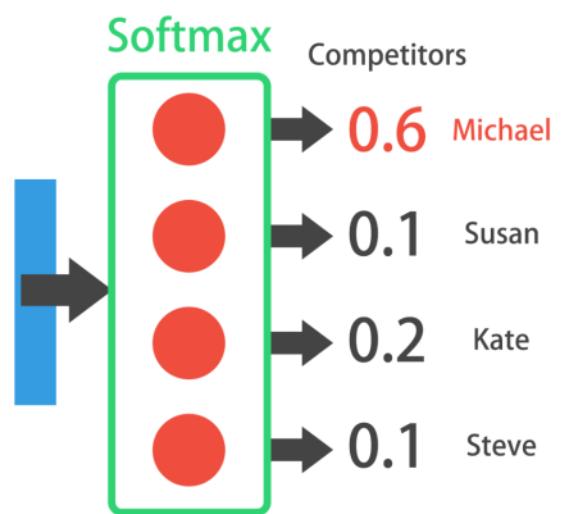
```
# Compile model  
model.compile(optimizer='sgd',  
              loss='binary_crossentropy')  
  
# Train model  
model.train(coordinates, labels, epochs=20)  
  
# Predict with trained model  
preds = model.predict(coordinates)
```



The architecture



The output layer



Multi-class model

```
# Instantiate a sequential model  
# ...  
  
# Add an input and hidden layer  
# ...  
  
# Add more hidden layers  
# ...  
  
# Add your output layer  
model.add(Dense(4, activation='softmax'))
```

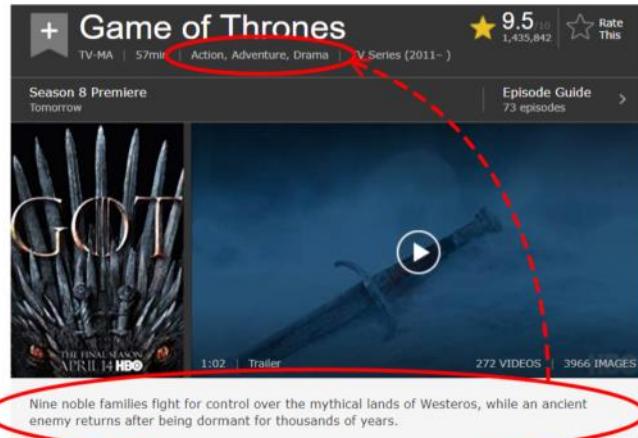
Preparing a dataset

```
import pandas as pd  
from keras.utils import to_categorical  
  
# Load dataset  
df = pd.read_csv('data.csv')  
  
# Turn response variable into labeled codes  
df.response = pd.Categorical(df.response)  
df.response = df.response.cat.codes  
  
# Turn response variable into one-hot response vector  
y = to_categorical(df.response)
```

One-hot encoding

Label Encoding			One Hot Encoding			
Food Name	Categorical #	Calories	Apple	Chicken	Broccoli	Calories
Apple	1	95	1	0	0	95
Chicken	2	231	0	1	0	231
Broccoli	3	50	0	0	1	50

Real world examples



	Multi-Class	Multi-Label
C = 3	Samples Labels (t) [0 0 1] [1 0 0] [0 1 0]	Samples Labels (t) [1 0 1] [0 1 0] [1 1 1]

However in a multi-label problem each individual in the sample can have all, none or a subset of the available classes.
<https://grombru.github.io/2018/05/23/multi-class-vs-multi-label.html>

The architecture

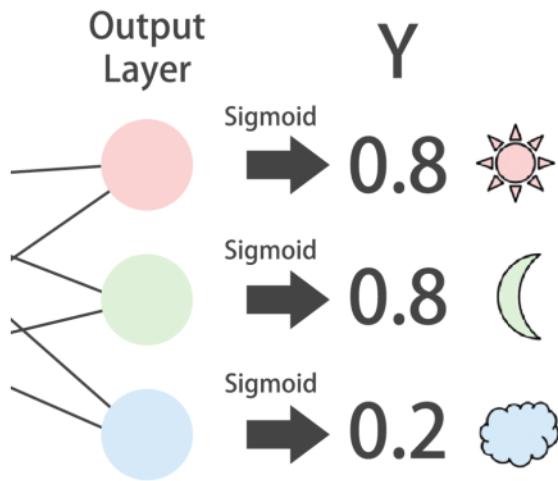
```
from keras.models import Sequential
from keras.layers import Dense

# Instantiate model
model = Sequential()

# Add input and hidden layers
model.add(Dense(2, input_shape=(1,)))

# Add an output layer for the 3 classes and sigmoid activation
model.add(Dense(3, activation='sigmoid'))
```

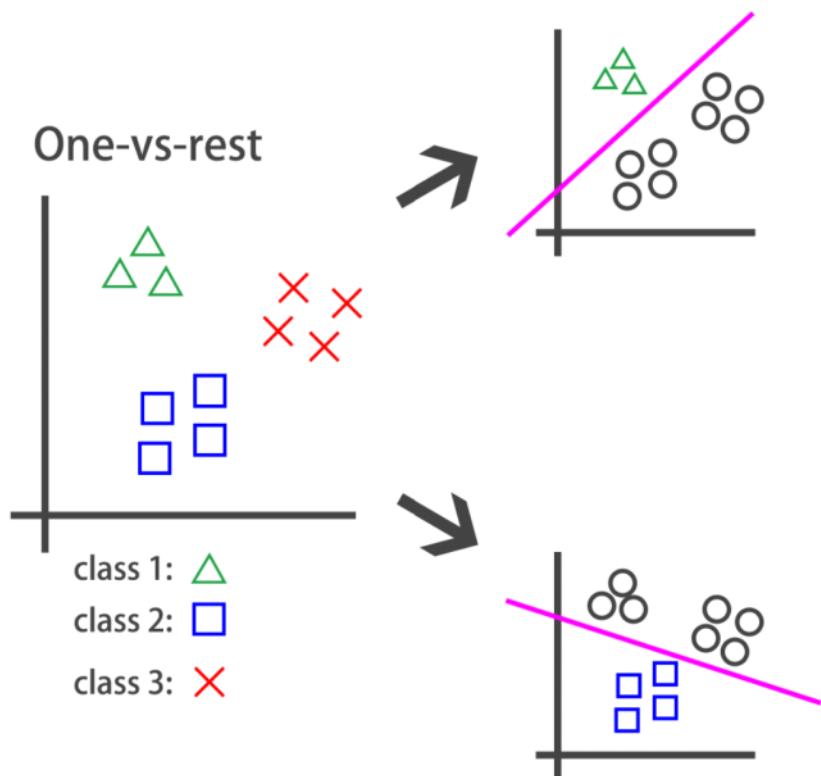
Sigmoid outputs



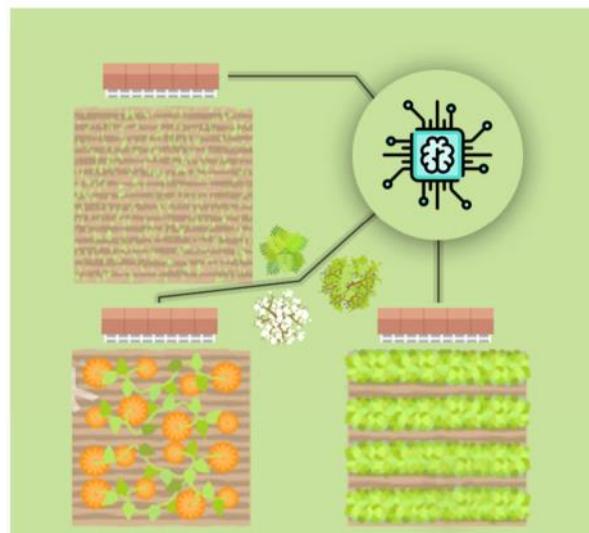
```
# Compile the model with binary crossentropy  
model.compile(optimizer='adam', loss='binary_crossentropy')
```

```
# Train your model, recall validation_split  
model.fit(X_train, y_train,  
          epochs=100,  
          validation_split=0.2)
```

```
Train on 1260 samples, validate on 280 samples  
Epoch 1/100  
1260/1260 [=====] - 0s 285us/step  
- loss: 0.7035 - acc: 0.6690 - val_loss: 0.5178 - val_acc: 0.7714  
...
```



An irrigation machine



An irrigation machine

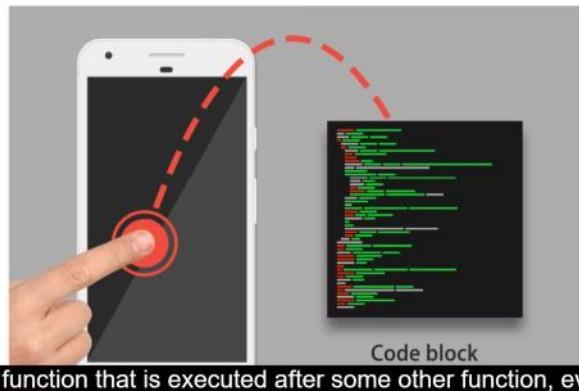
Sensor measurements

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	
0	4.0	3.0	4.0	4.0	1.0	1.0	4.0	0.0	3.0	3.0	2.0	2.0	3.0	1.0	4.0	2.0	2.0	2.0	1.0	4.0	
1	1.0	1.0	6.0	3.0	2.0	3.0	4.0	5.0	1.0	3.0	3.0	2.0	3.0	2.0	2.0	4.0	0.0	1.0	2.0	4.0	
2	1.0	5.0	7.0	6.0	4.0	0.0	0.0	6.0	0.0	1.0	1.0	3.0	4.0	2.0	1.0	0.0	4.0	1.0	1.0	3.0	
3	0.0	1.0	3.0	3.0	7.0	1.0	2.0	2.0	0.0	4.0	3.0	2.0	4.0	2.0	2.0	0.0	2.0	3.0	4.0	1.0	2.0
4	1.0	5.0	2.0	2.0	1.0	0.0	3.0	3.0	1.0	1.0	5.0	1.0	1.0	2.0	2.0	4.0	3.0	3.0	0.0	5.0	

Parcels to water

	0	1	2
0	0	0	0
1	1	1	1
2	1	1	0
3	1	1	1
4	0	0	0

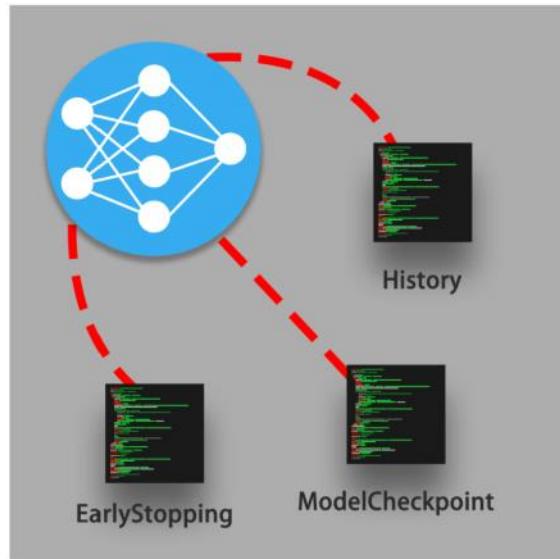
What is a callback?



A callback is a function that is executed after some other function, event, or task has finished.

For instance, when you touch your phone screen, a block of code that identifies the type of gesture will be triggered.

Callbacks in Keras



A callback you've been missing

```
# Training a model and saving its history
history = model.fit(X_train, y_train,
                     epochs=100,
                     metrics=[ 'accuracy' ])
print(history.history[ 'loss' ])
```

```
[0.6753975939750672, ..., 0.3155936544282096]
```

```
print(history.history[ 'acc' ])
```

```
[0.7030952412741525, ..., 0.8604761900220599]
```

A callback you've been missing

```
# Training a model and saving its history
history = model.fit(X_train, y_train,
                     epochs=100,
                     validation_data=(X_test, y_test),
                     metrics=['accuracy'])

print(history.history['val_loss'])
```

```
[0.7753975939750672, ..., 0.4155936544282096]
```

```
print(history.history['val_acc'])
```

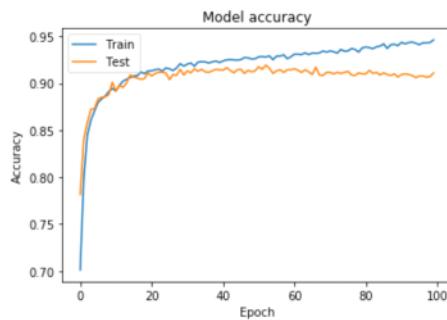
```
[0.6030952412741525, ..., 0.7604761900220599]
```

History plots

```
# Plot train vs test accuracy per epoch
plt.figure()

# Use the history metrics
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])

# Make it pretty
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'])
plt.show()
```



Early stopping

```
# Import early stopping from keras callbacks
from keras.callbacks import EarlyStopping

# Instantiate an early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train your model with the callback
model.fit(X_train, y_train, epochs=100,
           validation_data=(X_test, y_test),
           callbacks = [early_stopping])
```

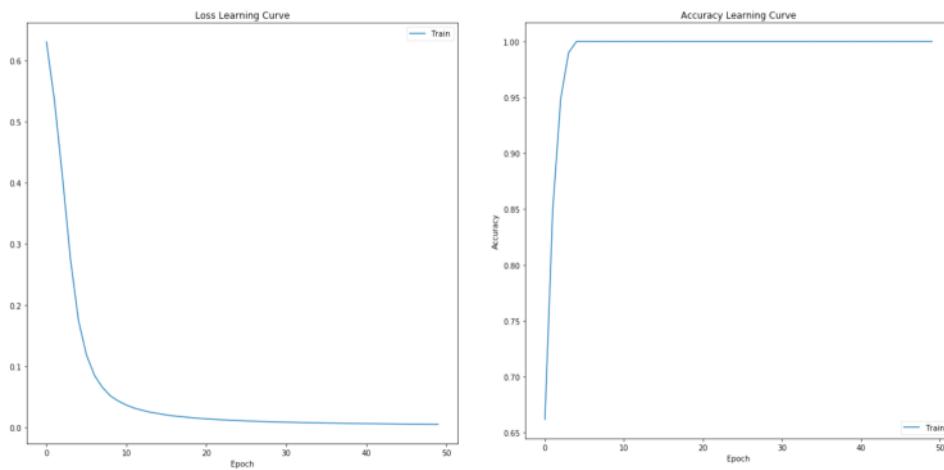
Model checkpoint

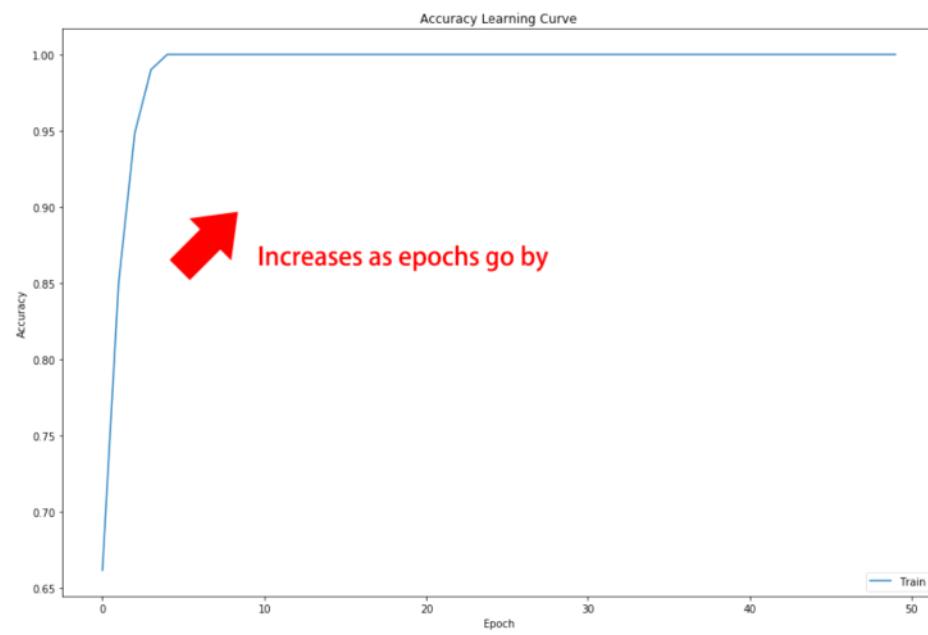
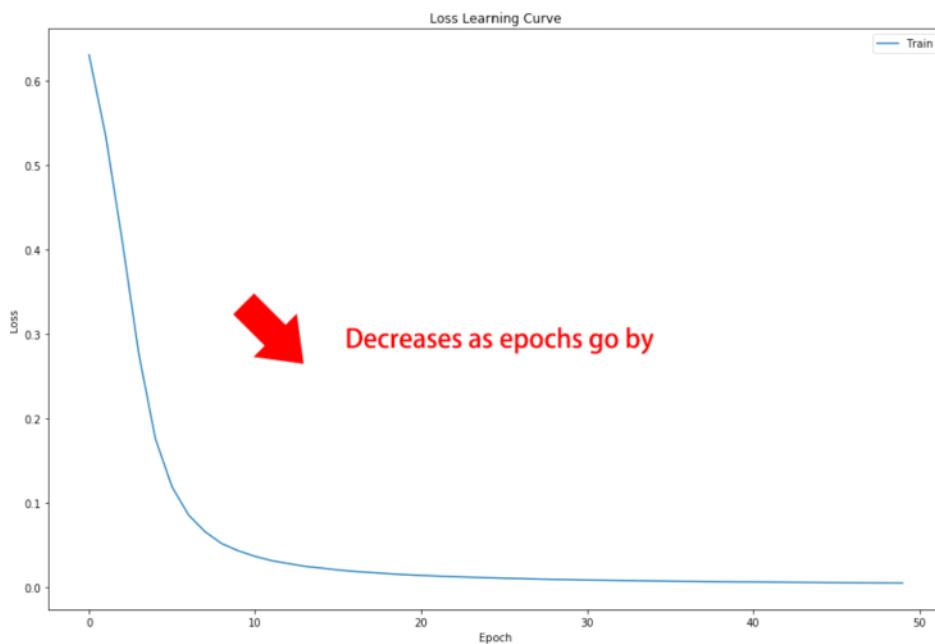
```
# Import model checkpoint from keras callbacks
from keras.callbacks import ModelCheckpoint

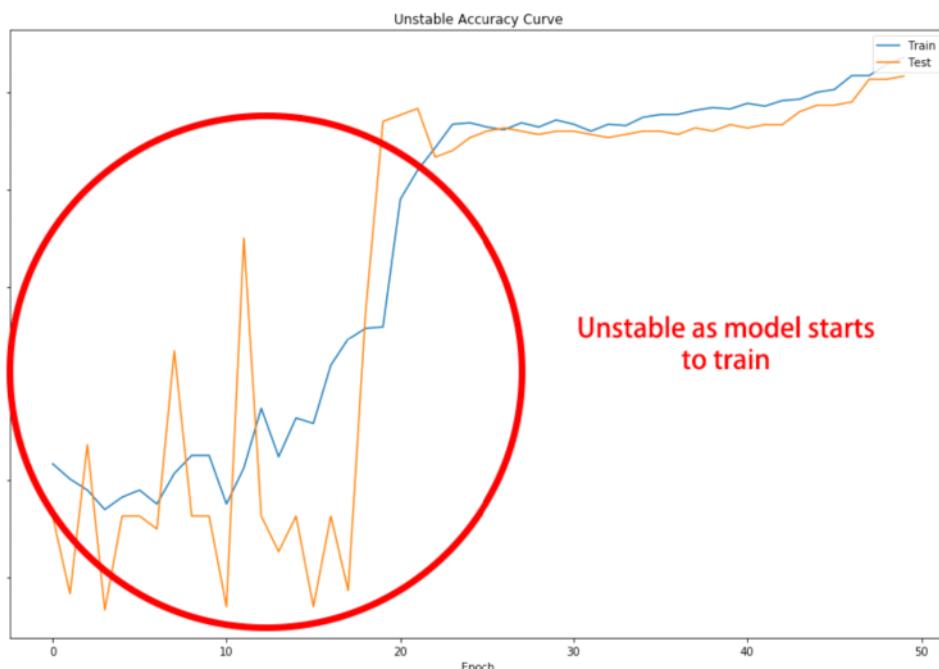
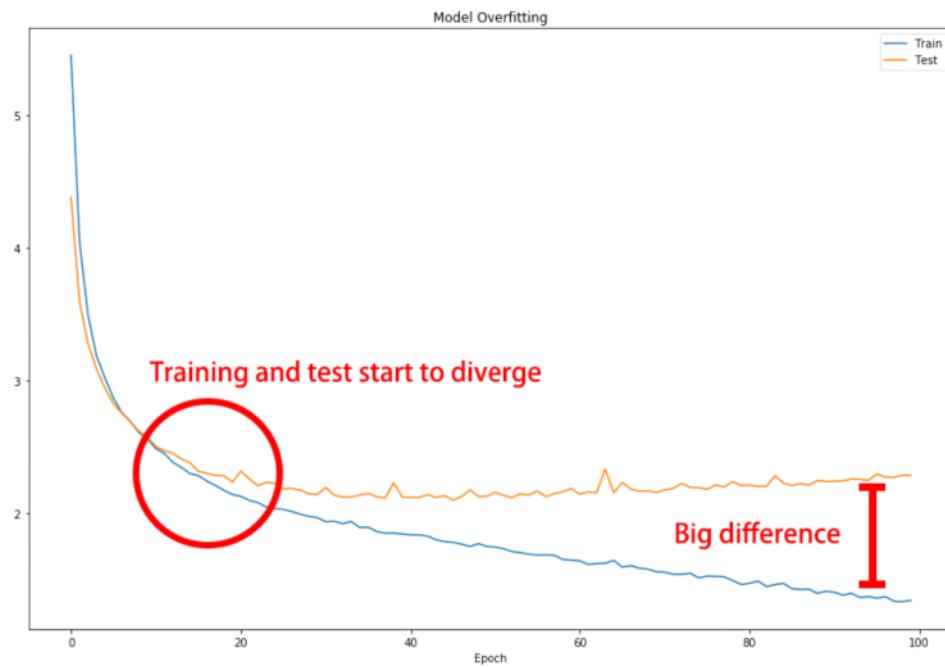
# Instantiate a model checkpoint callback
model_save = ModelCheckpoint('best_model.hdf5',
                             save_best_only=True)

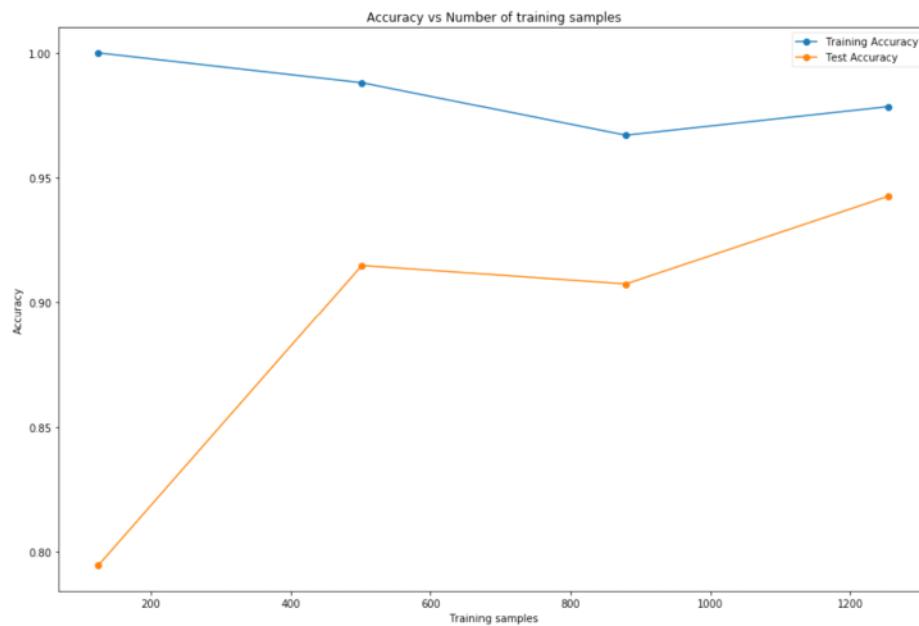
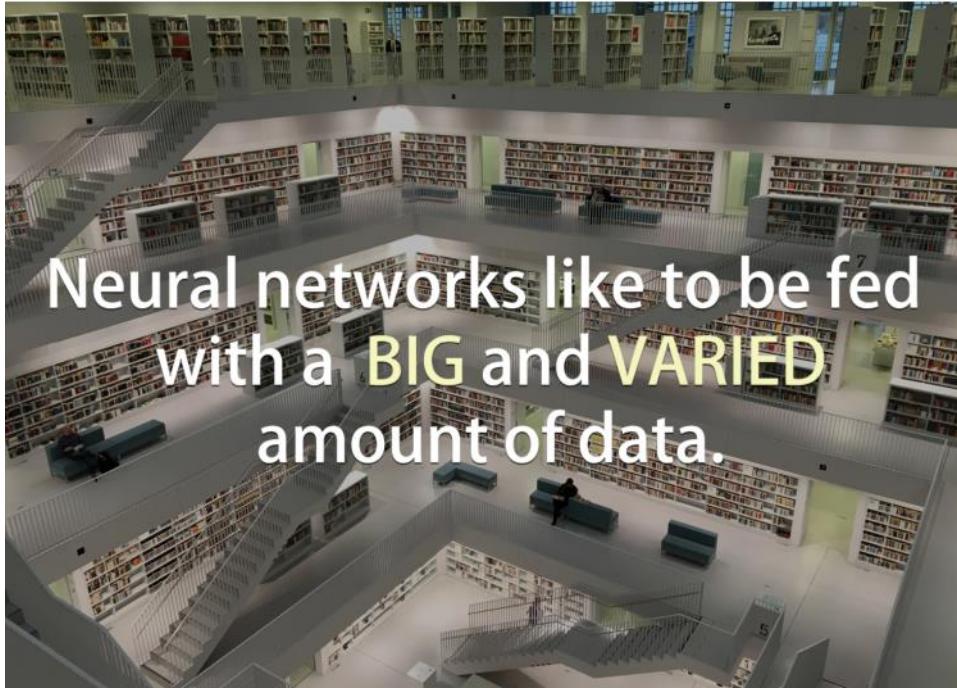
# Train your model with the callback
model.fit(X_train, y_train, epochs=100,
           validation_data=(X_test, y_test),
           callbacks = [model_save])
```

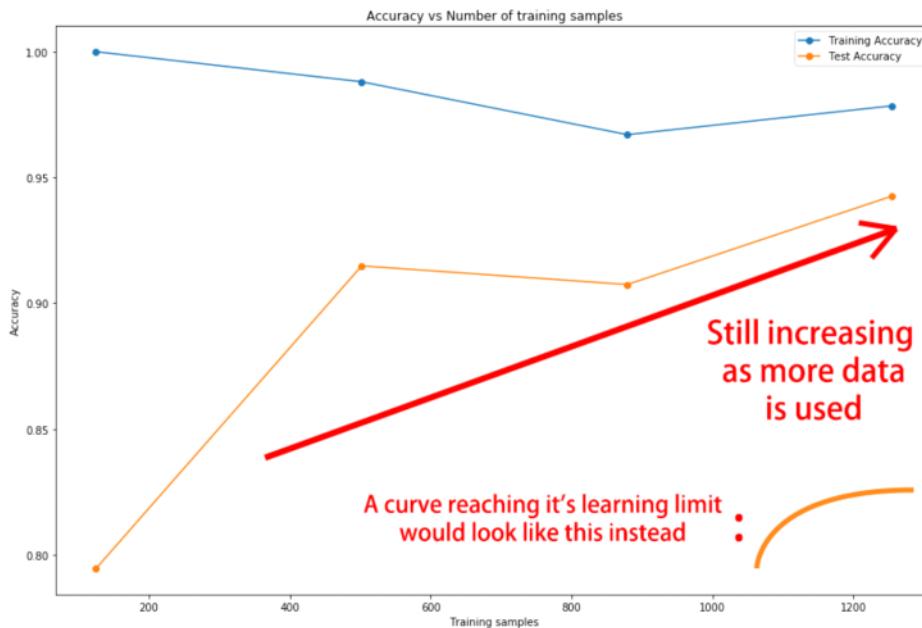
Improving Your Model Performance







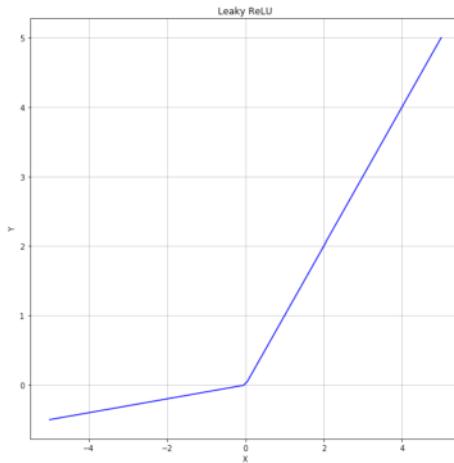
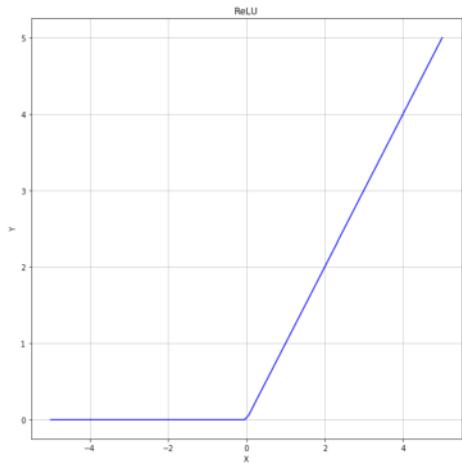
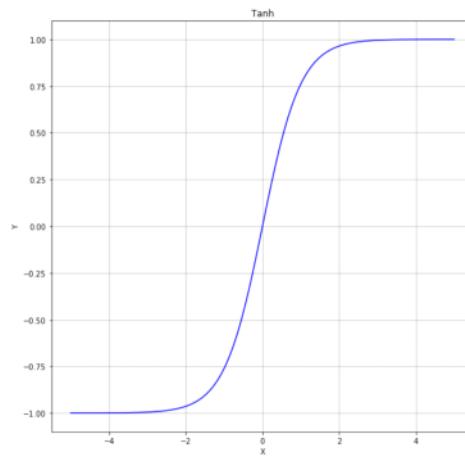
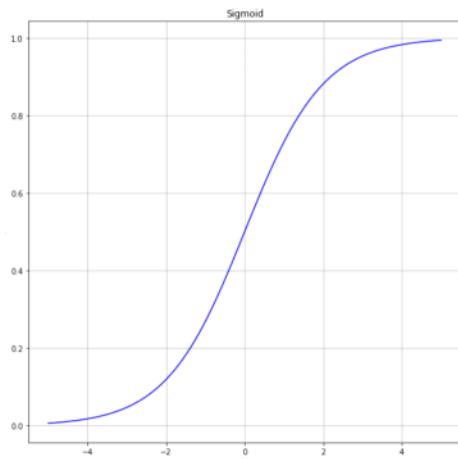
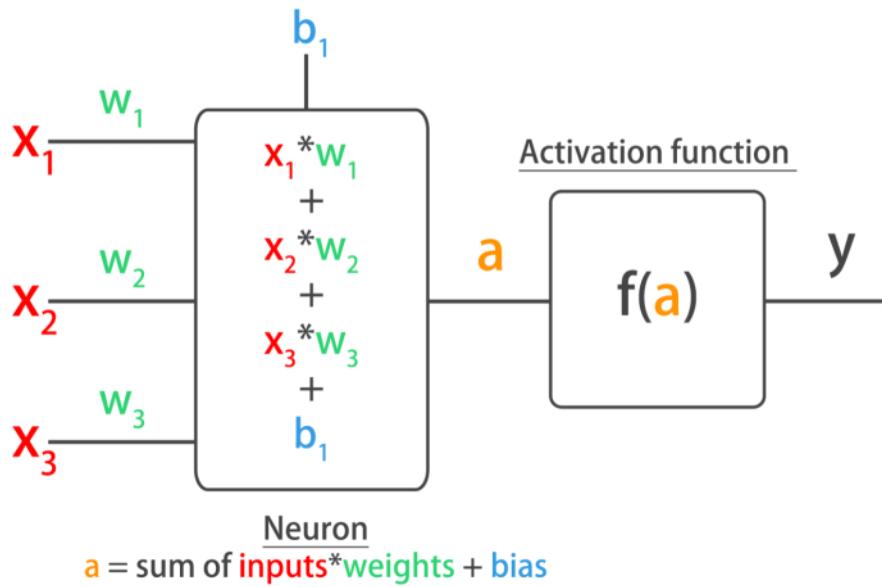




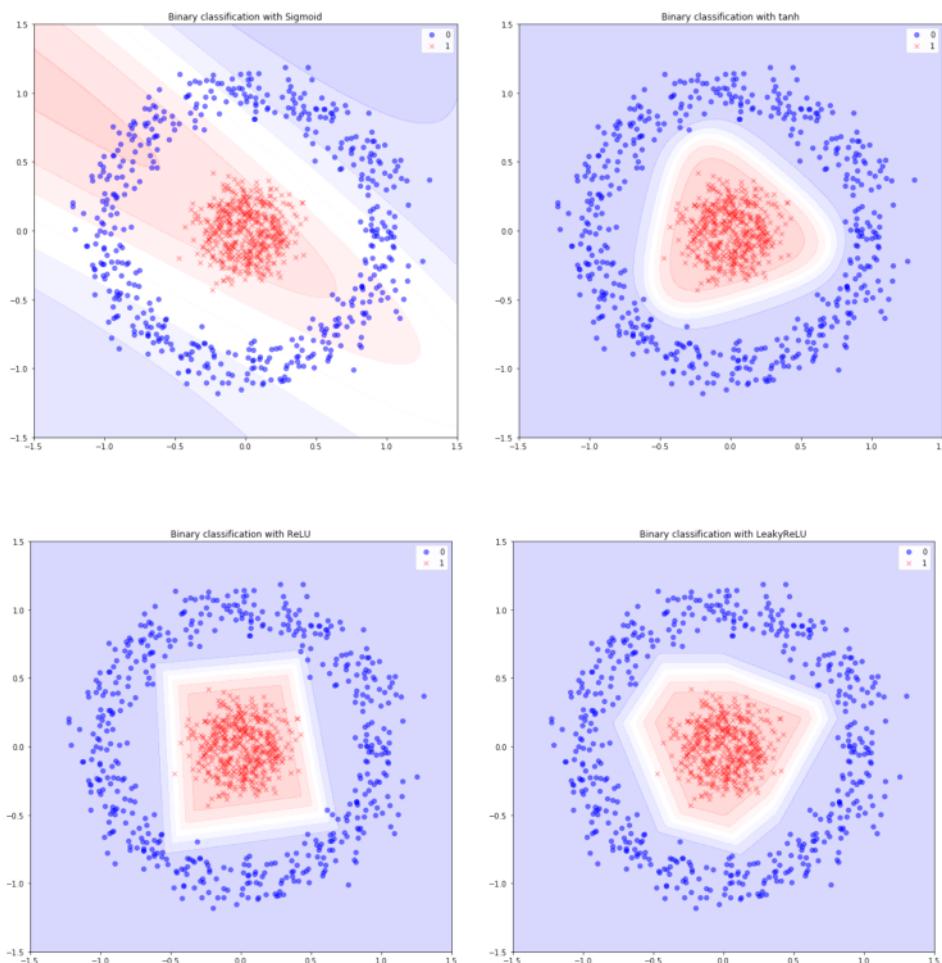
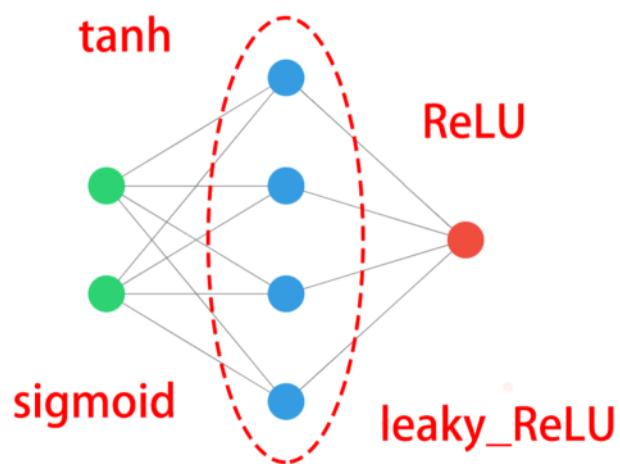
```
# Store initial model weights
init_weights = model.get_weights()

# Lists for storing accuracies
train_accs = []
tests_accs = []
```

```
for train_size in train_sizes:
    # Split a fraction according to train_size
    X_train_frac, _, y_train_frac, _ =
        train_test_split(X_train, y_train, train_size=train_size)
    # Set model initial weights
    model.set_weights(initial_weights)
    # Fit model on the training set fraction
    model.fit(X_train_frac, y_train_frac, epochs=100,
               verbose=0,
               callbacks=[EarlyStopping(monitor='loss', patience=1)])
    # Get the accuracy for this training set fraction
    train_acc = model.evaluate(X_train_frac, y_train_frac, verbose=0)[1]
    train_accs.append(train_acc)
    # Get the accuracy on the whole test set
    test_acc = model.evaluate(X_test, y_test, verbose=0)[1]
    test_accs.append(test_acc)
print("Done with size: ", train_size)
```

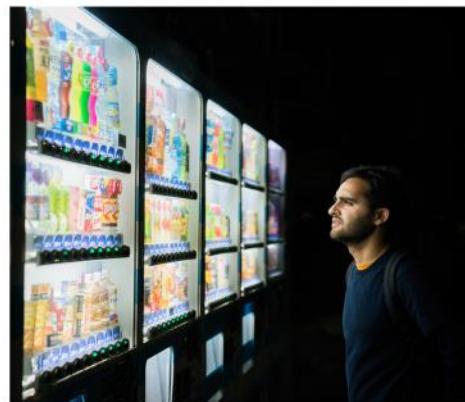


Effects of activation functions



Which activation function to use?

- No magic formula
- Different properties
- Depends on our problem
- Goal to achieve in a given layer
- ReLU are a good first choice
- Sigmoids not recommended for deep models
- Tune with experimentation



Comparing activation functions

```
# Set a random seed
np.random.seed(1)

# Return a new model with the given activation
def get_model(act_function):
    model = Sequential()
    model.add(Dense(4, input_shape=(2,), activation=act_function))
    model.add(Dense(1, activation='sigmoid'))
    return model
```

Comparing activation functions

```
# Activation functions to try out
activations = [ 'relu', 'sigmoid', 'tanh' ]

# Dictionary to store results
activation_results = {}

for funct in activations:
    model = get_model(act_function=funct)
    history = model.fit(X_train, y_train,
                         validation_data=(X_test, y_test),
                         epochs=100, verbose=0)
    activation_results[funct] = history
```

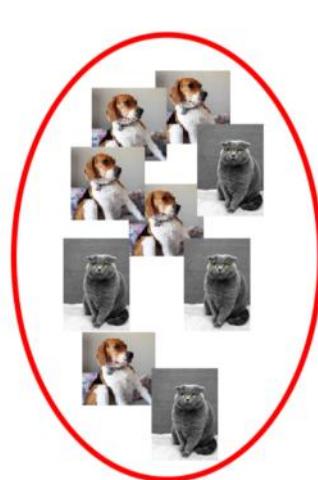
Comparing activation functions

```
import pandas as pd

# Extract val_loss history of each activation function
val_loss_per_funct = {k:v.history['val_loss'] for k,v in activation_results.items()}

# Turn the dictionary into a pandas dataframe
val_loss_curves = pd.DataFrame(val_loss_per_funct)

# Plot the curves
val_loss_curves.plot(title='Loss per Activation function')
```

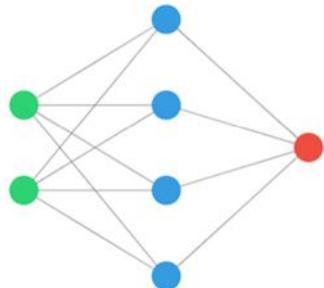


Batch

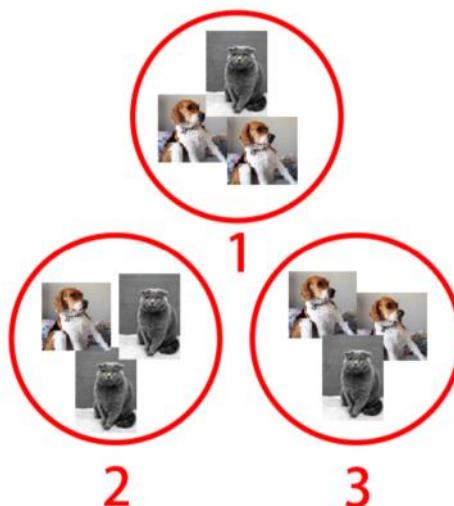


Mini-batches

The network is fed with 3 mini-batches



1 Epoch = 3 weight updates



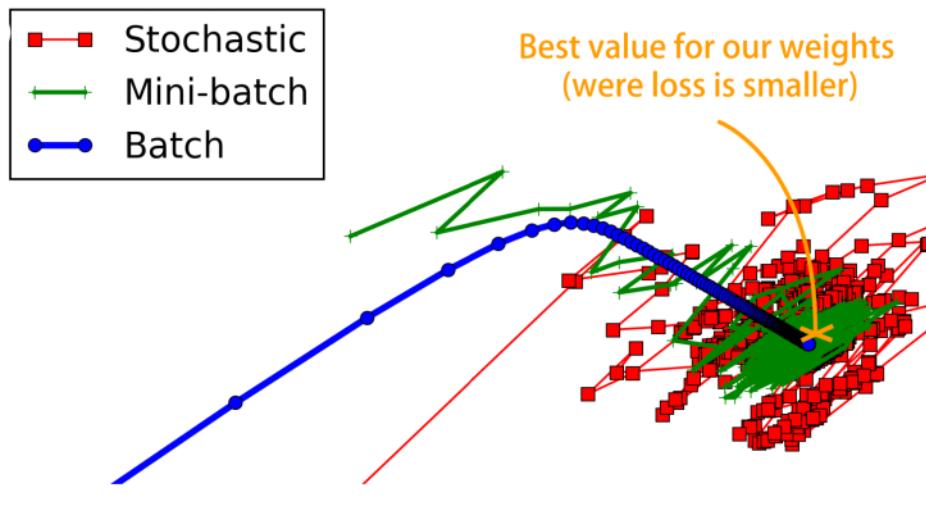
Mini-batches

Advantages

- Networks train faster (more weight updates in same amount of time)
- Less RAM memory required, can train on huge datasets
- Noise can help networks reach a lower error, escaping local minima

Disadvantages

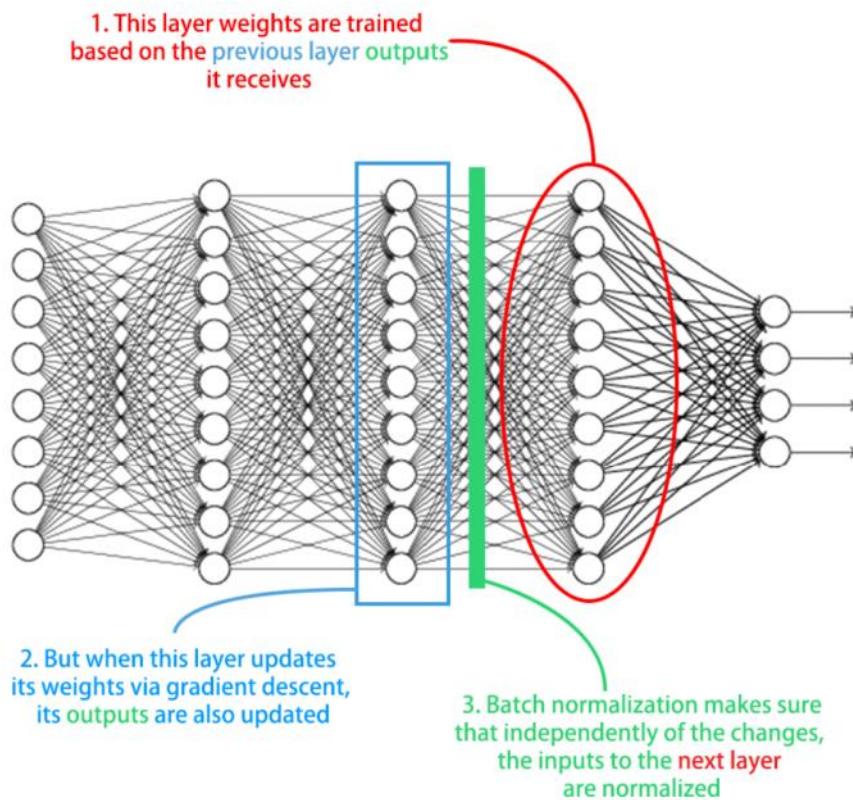
- More iterations need to be run
- Need to be adjusted, we need to find a good batch size



Standardization
(a normalization approach)

data - mean

standard deviation



Batch normalization advantages

- Improves gradient flow
- Allows higher learning rates
- Reduces dependence on weight initializations
- Acts as an unintended form of regularization
- Limits internal covariate shift

Batch normalization in Keras

```
# Import BatchNormalization from keras layers
from keras.layers import BatchNormalization

# Instantiate a Sequential model
model = Sequential()

# Add an input layer
model.add(Dense(3, input_shape=(2,), activation = 'relu'))

# Add batch normalization for the outputs of the layer above
model.add(BatchNormalization())

# Add an output layer
model.add(Dense(1, activation='sigmoid'))
```

Neural network hyperparameters

- Number of layers
- Number of neurons per layer
- Layer order
- Layer activations
- Batch sizes
- Learning rates
- Optimizers

Sklearn recap

```
# Import RandomizedSearchCV
from sklearn.model_selection import RandomizedSearchCV
# Instantiate your classifier
tree = DecisionTreeClassifier()
# Define a series of parameters to look over
params = {'max_depth':[3,None], "max_features":range(1,4), 'min_samples_leaf': range(1,4)}
# Perform random search with cross validation
tree_cv = RandomizedSearchCV(tree, params, cv=5)
tree_cv.fit(X,y)

# Print the best parameters
print(tree_cv.best_params_)
```

```
{'min_samples_leaf': 1, 'max_features': 3, 'max_depth': 3}
```

Turn a Keras model into a Sklearn estimator

```
# Function that creates our Keras model
def create_model(optimizer='adam', activation='relu'):
    model = Sequential()
    model.add(Dense(16, input_shape=(2,), activation=activation))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer=optimizer, loss='binary_crossentropy')
    return model

# Import sklearn wrapper from keras
from keras.wrappers.scikit_learn import KerasClassifier

# Create a model as a sklearn estimator
model = KerasClassifier(build_fn=create_model, epochs=6, batch_size=16)
```

Cross-validation

```
# Import cross_val_score
from sklearn.model_selection import cross_val_score

# Check how your keras model performs with 5 fold crossvalidation
kfold = cross_val_score(model, X, y, cv=5)

# Print the mean accuracy per fold
kfold.mean()
```

```
0.913333
```

```
# Print the standard deviation per fold
kfold.std()
```

Tips for neural networks hyperparameter tuning

- Random search is preferred over grid search
- Don't use many epochs
- Use a smaller sample of your dataset
- Play with batch sizes, activations, optimizers and learning rates

Random search on Keras models

```
# Define a series of parameters
params = dict(optimizer=['sgd', 'adam'], epochs=3,
               batch_size=[5, 10, 20], activation=['relu', 'tanh'])

# Create a random search cv object and fit it to the data
random_search = RandomizedSearchCV(model, params_dist=params, cv=3)
random_search_results = random_search.fit(X, y)

# Print results
print("Best: %f using %s".format(random_search_results.best_score_,
random_search_results.best_params_))

Best: 0.94 using {'optimizer': 'adam', 'epochs': 3, 'batch_size': 10, 'activation': 'relu'}
```

Tuning other hyperparameters

```
def create_model(nl=1,nn=256):
    model = Sequential()
    model.add(Dense(16, input_shape=(2,), activation='relu'))

    # Add as many hidden layers as specified in nl
    for i in range(nl):
        # Layers have nn neurons
        model.add(Dense(nn, activation='relu'))

    # End defining and compiling your model...
```

Tuning other hyperparameters

```
# Define parameters, named just like in create_model()
params = dict(nl=[1, 2, 9], nn=[128, 256, 1000])

# Repeat the random search...

# Print results...
```

```
Best: 0.87 using {'nl': 2, 'nn': 128}
```

Advance Model Architecture

Accessing Keras layers

```
# Acessing the first layer of a Keras model
first_layer = model.layers[0]

# Printing the layer, and its input, output and weights
print(first_layer.input)
print(first_layer.output)
print(first_layer.weights)
```

```
<tf.Tensor 'dense_1_input:0' shape=(?, 3) dtype=float32>
```

```
<tf.Tensor 'dense_1/Relu:0' shape=(?, 2) dtype=float32>
```

```
[<tf.Variable 'dense_1/kernel:0' shape=(3, 2) dtype=float32_ref>,
 <tf.Variable 'dense_1/bias:0' shape=(2,) dtype=float32_ref>]
```

What are tensors?

```
# Defining a rank 2 tensor (2 dimensions)
T2 = [[1,2,3],
      [4,5,6],
      [7,8,9]]

# Defining a rank 3 tensor (3 dimensions)
T3 = [[1,2,3],
      [4,5,6],
      [7,8,9],

      [10,11,12],
      [13,14,15],
      [16,17,18],

      [19,20,21],
      [22,23,24],
      [25,26,27]]
```

```
# Import Keras backend
import keras.backend as K

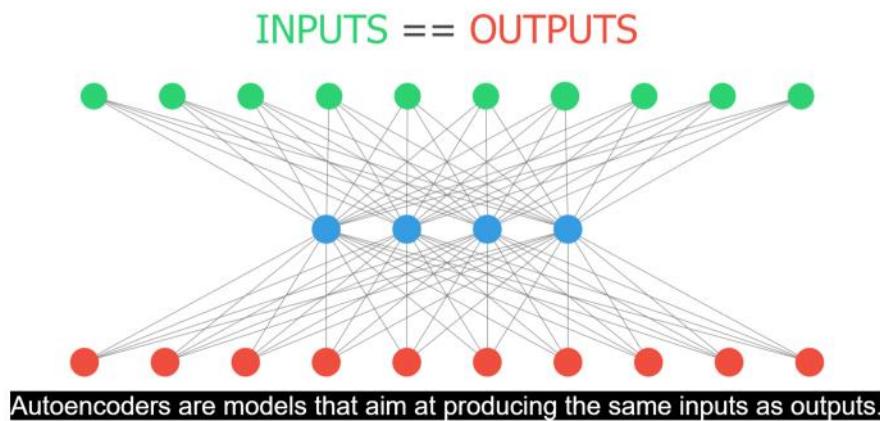
# Get the input and output tensors of a model layer
inp = model.layers[0].input
out = model.layers[0].output

# Function that maps layer inputs to outputs
inp_to_out = K.function([inp], [out])

# We pass an input and get the output we'd get in that first layer
print(inp_to_out([X_train]))
```

```
# Outputs of the first layer per sample in X_train
[array([[0.7, 0], ..., [0.1, 0.3]])]
```

Autoencoders!



Autoencoder use cases

- Dimensionality reduction:
 - Smaller dimensional space representation of our inputs.
- De-noising data:
 - If trained with clean data, irrelevant noise will be filtered out during reconstruction.
- Anomaly detection:
 - A poor reconstruction will result when the model is fed with unseen inputs.



Building a simple autoencoder

```
# Instantiate a sequential model
autoencoder = Sequential()

# Add a hidden layer of 4 neurons and an input layer of 100
autoencoder.add(Dense(4, input_shape=(100,), activation='relu'))

# Add an output layer of 100 neurons
autoencoder.add(Dense(100, activation='sigmoid'))

# Compile your model with the appropriate loss
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

Breaking it into an encoder

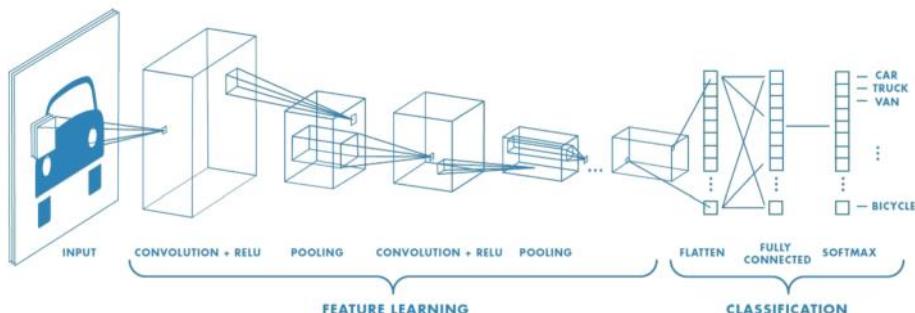
```
# Building a separate model to encode inputs
encoder = Sequential()
encoder.add(autoencoder.layers[0])

# Predicting returns the four hidden layer neuron outputs
encoder.predict(X_test)
```

```
# Four numbers for each observation in X_test
array([10.0234375, 5.833543, 18.90444, 9.20348],...)
```

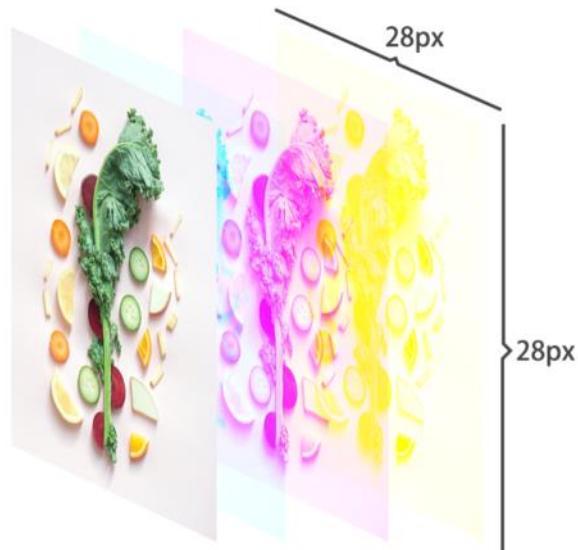
3	3	2	1	0
0	0	1 ₀	3 ₁	1 ₂
3	1	2 ₂	2 ₂	3 ₀
2	0	0 ₀	2 ₁	2 ₂
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



```
    input_shape  
(WIDTH,HEIGHT,CHANNELS)  
    input_shape = (28,28,3)
```

Color Images have 3 channels
RGB (Red Green Blue)



```
# Import Conv2D layer and Flatten from keras layers  
from keras.layers import Dense, Conv2D, Flatten  
  
# Instantiate your model as usual  
model = Sequential()  
  
# Add a convolutional layer with 32 filters of size 3x3  
model.add(Conv2D(filters=32,  
                 kernel_size=3,  
                 input_shape=(28, 28, 1),  
                 activation='relu'))  
  
# Add another convolutional layer  
model.add(Conv2D(8, kernel_size=3, activation='relu'))  
  
# Flatten the output of the previous layer  
model.add(Flatten())  
  
# End this multiclass model with 3 outputs and softmax  
model.add(Dense(3, activation='softmax'))
```

Pre-processing images for ResNet50

```
# Import image from keras preprocessing
from keras.preprocessing import image

# Import preprocess_input from keras applications resnet50
from keras.applications.resnet50 import preprocess_input

# Load the image with the right target size for your model
img = image.load_img(img_path, target_size=(224, 224))

# Turn it into an array
img = image.img_to_array(img)

# Expand the dimensions so that it's understood by our network:
# img.shape turns from (224, 224, 3) into (1, 224, 224, 3)
img = np.expand_dims(img, axis=0)

# Pre-process the img in the same way training images were
img = preprocess_input(img)
```

Using the ResNet50 model in Keras

```
# Import ResNet50 and decode_predictions from keras.applications.resnet50
from keras.applications.resnet50 import ResNet50, decode_predictions

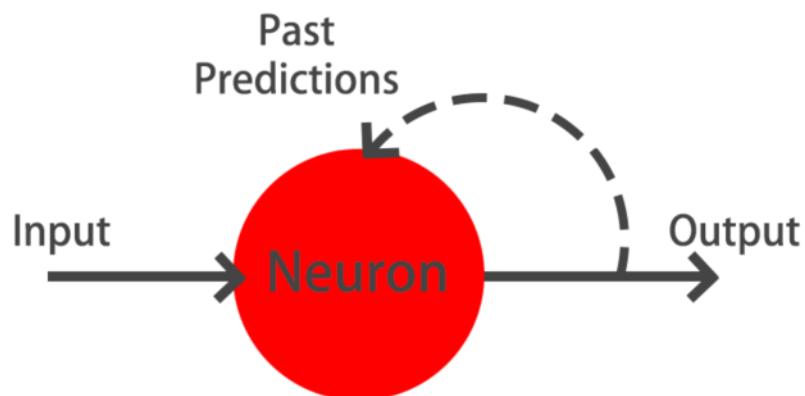
# Instantiate a ResNet50 model with imagenet weights
model = ResNet50(weights='imagenet')

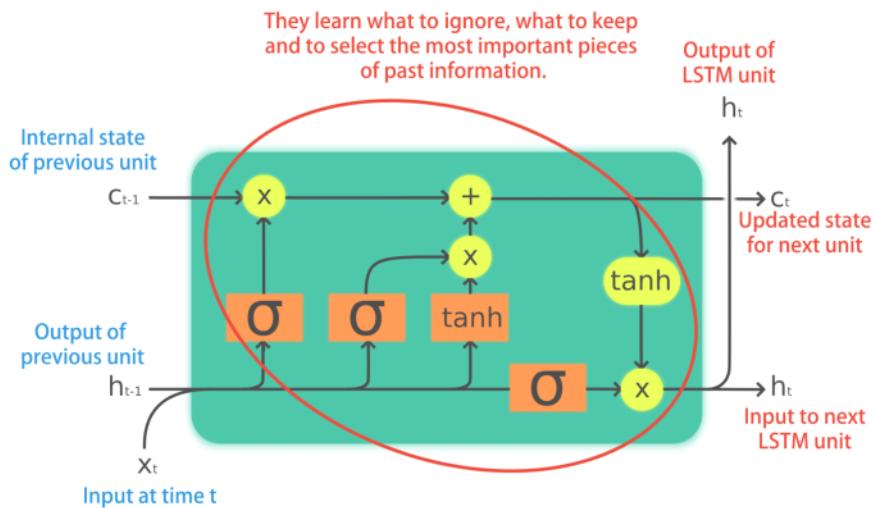
# Predict with ResNet50 on our img
preds = model.predict(img)

# Decode predictions and print it
print('Predicted:', decode_predictions(preds, top=1)[0])
```

```
Predicted: [('n07697313', 'cheeseburger', 0.9868016)]
```

What are RNNs?





When to use LSTMs?

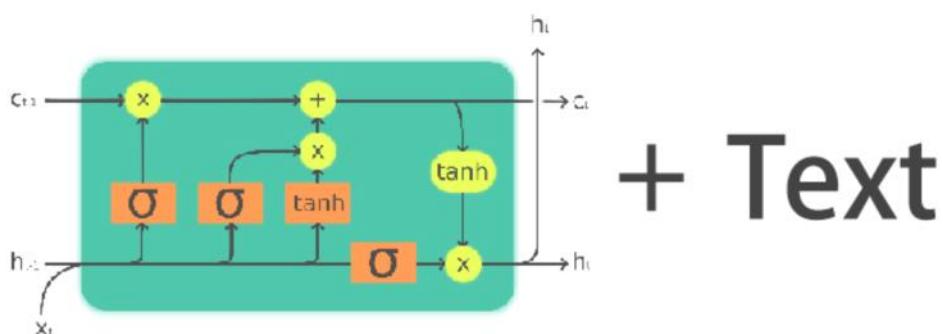
- Image captioning
- Speech to text
- Text translation
- Document summarization
- Text generation
- Musical composition
- ...



'man in black shirt is playing guitar.'



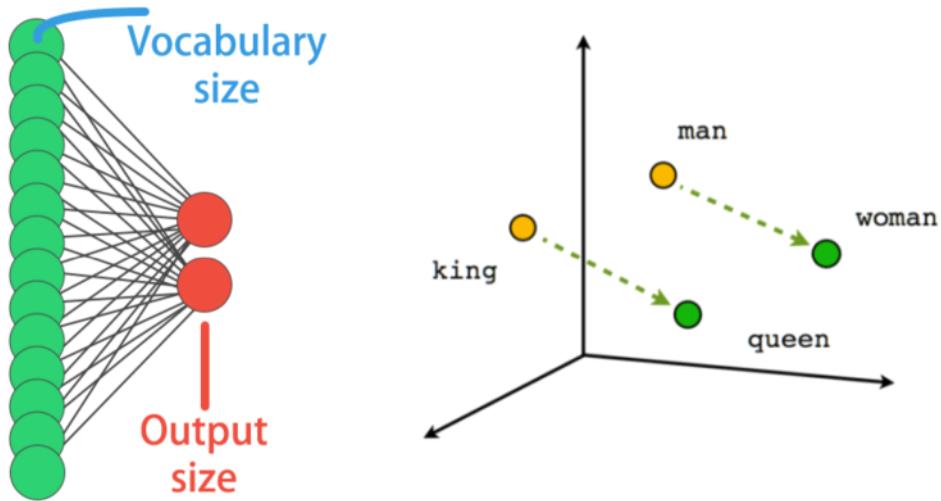
'construction worker in orange safety vest is working on road.'



this is a sentence



42 11 23 1



```
text = 'Hi this is a small sentence'
```

```
# We choose a sequence length
```

```
seq_len = 3
```

```
# Split text into a list of words
```

```
words = text.split()
```

```
['Hi', 'this', 'is', 'a', 'small', 'sentence']
```

```
# Make lines
```

```
lines = []
```

```
for i in range(seq_len, len(words) + 1):
```

```
    line = ' '.join(words[i-seq_len:i])
```

```
    lines.append(line)
```

```
['Hi this is', 'this is a', 'is a small', 'a small sentence']
```

```
# Import Tokenizer from keras preprocessing text
```

```
from keras.preprocessing.text import Tokenizer
```

```
# Instantiate Tokenizer
```

```
tokenizer = Tokenizer()
```

```
# Fit it on the previous lines
```

```
tokenizer.fit_on_texts(lines)
```

```
# Turn the lines into numeric sequences
```

```
sequences = tokenizer.texts_to_sequences(lines)
```

```
array([[5, 3, 1], [3, 1, 2], [1, 2, 4], [2, 4, 6]])
```

```
print(tokenizer.index_word)
```

```
{1: 'is', 2: 'a', 3: 'this', 4: 'small', 5: 'hi', 6: 'sentence'}
```

```
# Import Dense, LSTM and Embedding layers
from keras.layers import Dense, LSTM, Embedding
model = Sequential()

# Vocabulary size
vocab_size = len(tokenizer.index_word) + 1

# Starting with an embedding layer
model.add(Embedding(input_dim=vocab_size, output_dim=8, input_length=2))

# Adding an LSTM layer
model.add(LSTM(8))

# Adding a Dense hidden layer
model.add(Dense(8, activation='relu'))

# Adding an output layer with softmax
model.add(Dense(vocab_size, activation='softmax'))
```