

Image Processing with Keras in Python

Saturday, 22 August 2020 9:35 PM

Image Processing With Neural Networks

Images as data

```
import matplotlib.pyplot as plt  
data = plt.imread('stop_sign.jpg')  
plt.imshow(data)  
plt.show()
```



Images as data

```
data.shape
```

```
(2832, 4256, 3)
```

Images as data

```
data[1000, 1500]
```

```
array([0.73333333, 0.07843137, 0.14509804])
```



Images as data

```
data[250, 3500]
```

```
array([0.25882353, 0.43921569, 0.77254902])
```



Modifying image data

```
data[:, :, 1] = 0  
data[:, :, 2] = 0  
plt.imshow(data)  
plt.show()
```



For example, here we set the green and blue values of the pixels to zero.

Changing an image

```
data[200:1200, 200:1200, :] = [0, 1, 0]  
plt.imshow(data)  
plt.show()
```



Black and white images

```
tshirt[10:20, 15:25] = 1  
plt.imshow(tshirt)  
plt.show()
```

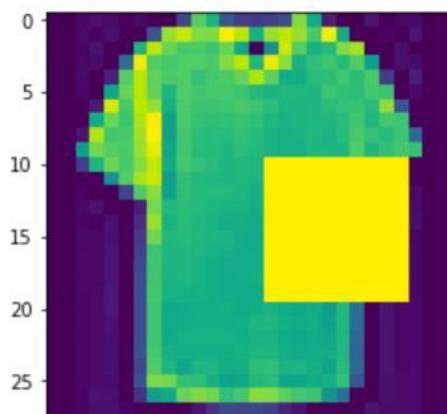


Image classification: training

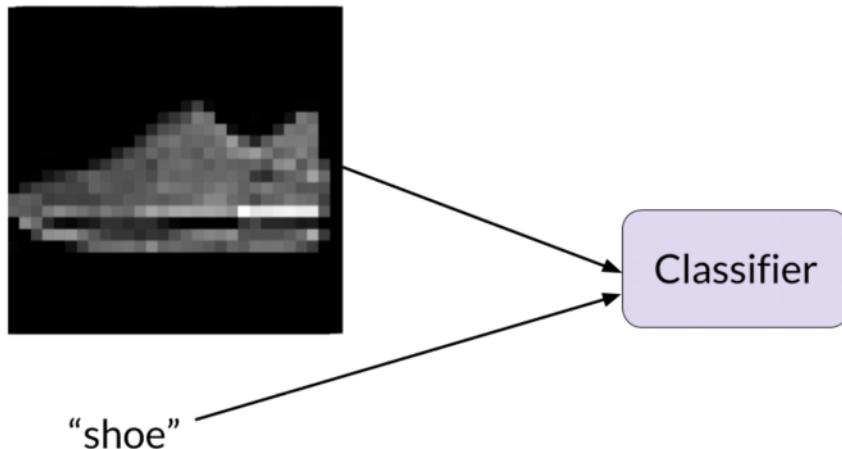
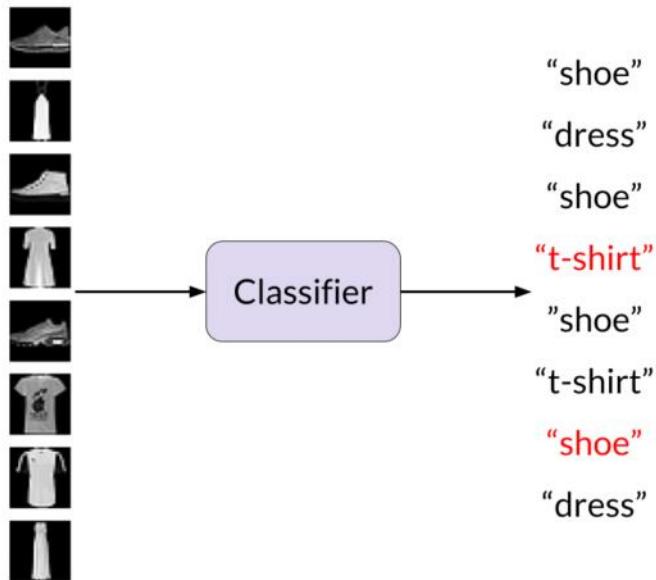


Image classification: evaluation



Representing class data: one-hot encoding

```
labels = ["shoe", "dress", "shoe", "t-shirt",
          "shoe", "t-shirt", "shoe", "dress"]
```

Representing class data: one-hot encoding

```
array([[0., 0., 1.],    <= shoe
       [0., 1., 0.],    <= dress
       [0., 0., 1.],    <= shoe
       [1., 0., 0.],    <= t-shirt
       [0., 0., 1.],    <= shoe
       [1., 0., 0.],    <= t-shirt
       [0., 0., 1.],    <= shoe
       [0., 1., 0.]])   <= dress
```



One-hot encoding

```
categories = np.array(["t-shirt", "dress", "shoe"])
n_categories = 3
ohe_labels = np.zeros((len(labels), n_categories))

for ii in range(len(labels)):
    jj = np.where(categories == labels[ii])
    ohe_labels[ii, jj] = 1
```

One-hot encoding: testing predictions

test	prediction
array([[0., 0., 1.], [0., 1., 0.], [0., 0., 1.], [0., 1., 0.], [0., 0., 1.], [0., 0., 1.], [0., 0., 1.], [0., 1., 0.]])	array([[0., 0., 1.], [0., 1., 0.], [0., 0., 1.], [1., 0., 0.], <= incorrect [0., 0., 1.], [1., 0., 0.], <= incorrect [0., 0., 1.], [0., 1., 0.]])

```
(test * prediction).sum()
```

6.0

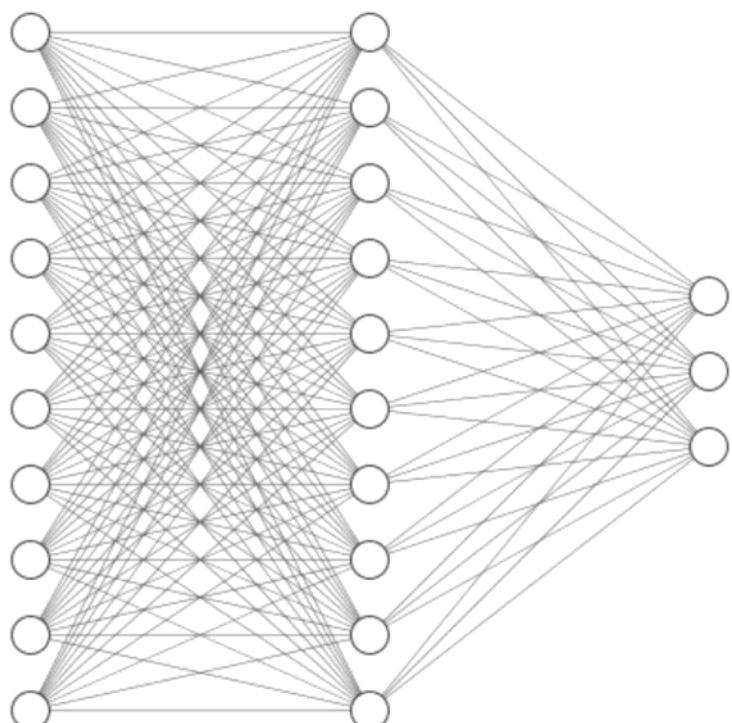
Keras for image classification

```
from keras.layers import Dense  
train_data.shape
```

(50, 28, 28, 1)

Keras for image classification

```
model.add(Dense(10, activation='relu',  
               input_shape=(784,)))  
  
model.add(Dense(10, activation='relu'))  
  
model.add(Dense(3, activation='softmax'))
```



Keras for image classification

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

Keras for image classification

```
train_data = train_data.reshape((50, 784))
```

Keras for image classification

```
model.fit(train_data, train_labels,  
          validation_split=0.2,  
          epochs=3)
```

```
model.fit(train_data, train_labels,  
          validation_split=0.2,  
          epochs=3)
```

```
Train on 40 samples, validate on 10 samples  
Epoch 1/3  
32/40 [=====>.....] - ETA: 0s - loss: 1.0117 - acc: 0.4688  
40/40 [=====] - 0s 4ms/step - loss: 1.0438 - acc: 0.4250  
                         - val_loss: 0.9668 - val_acc: 0.4000  
Epoch 2/3  
32/40 [=====>.....] - ETA: 0s - loss: 0.9556 - acc: 0.5312  
40/40 [=====] - 0s 195us/step - loss: 0.9404 - acc: 0.5750  
                         - val_loss: 0.9068 - val_acc: 0.4000  
Epoch 3/3  
32/40 [=====>.....] - ETA: 0s - loss: 0.9143 - acc: 0.5938  
40/40 [=====] - 0s 189us/step - loss: 0.8726 - acc: 0.6750  
                         - val_loss: 0.8452 - val_acc: 0.4000
```

Keras for image classification

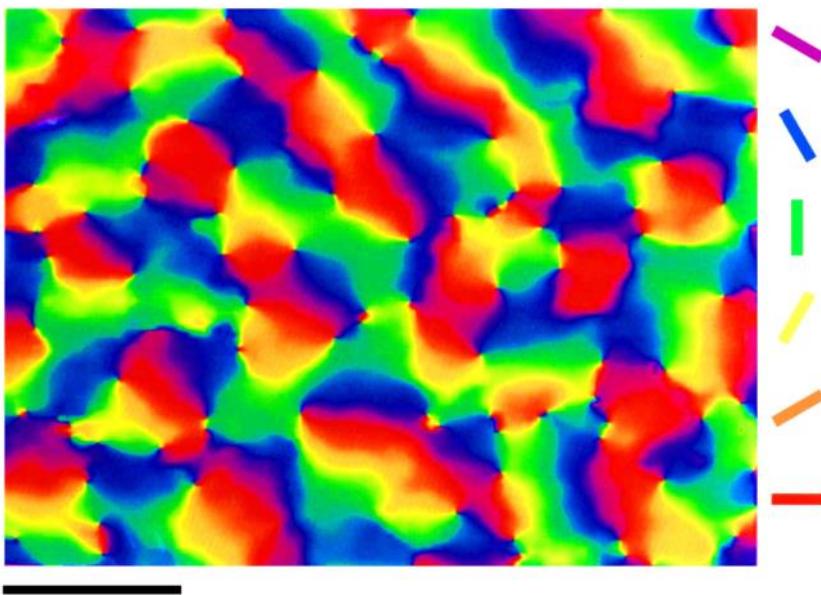
```
test_data = test_data.reshape((10, 784))  
model.evaluate(test_data, test_labels)
```

Using Convolutions

Using correlations in images

- Natural images contain spatial correlations
- For example, pixels along a contour or edge
- How can we use these correlations?

Biological inspiration



What is a convolution?

```
array = np.array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
kernel = np.array([-1, 1])
conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
conv[0] = (kernel * array[0:2]).sum()
conv[1] = (kernel * array[1:3]).sum()
conv[2] = (kernel * array[2:4]).sum()
...
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()
conv
```

Convolution in one dimension

```
array = np.array([0, 0, 1, 1, 0, 0, 1, 1, 0, 0])
kernel = np.array([-1, 1])
conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()
conv
```

```
array([ 0,  1,  0, -1,  0,  1,  0, -1,  0])
```

Image convolution

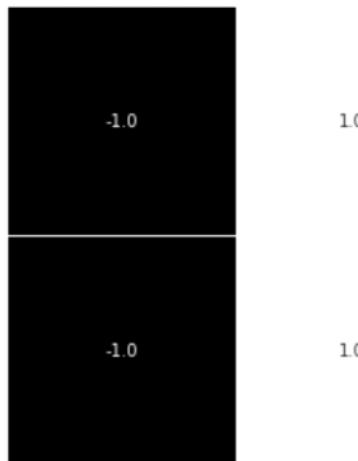
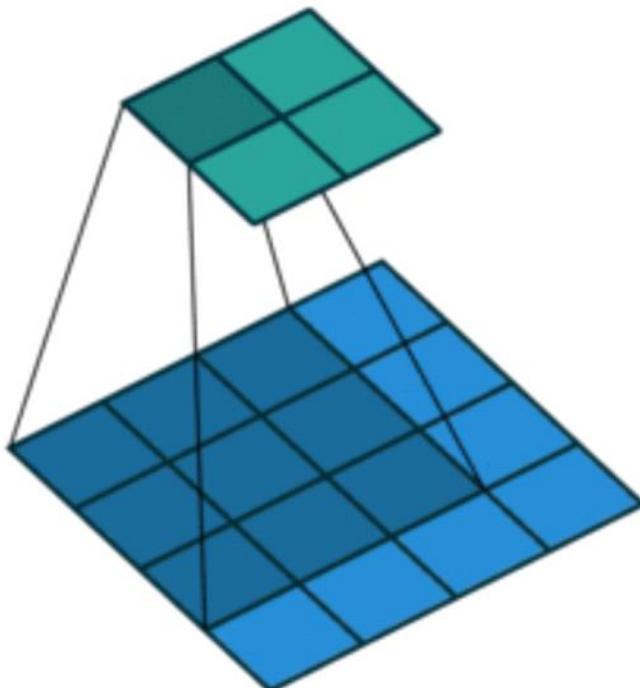


Image convolution

Two-dimensional convolution

```
kernel = np.array([[-1, 1],  
                  [-1, 1]])  
  
conv = np.zeros((27, 27))  
  
for ii in range(27):  
    for jj in range(27):  
        window = image[ii:ii+2, jj:jj+2]  
        conv[ii, jj] = np.sum(window * kernel)
```

Convolution



Keras Convolution layer

```
from keras.layers import Conv2D  
Conv2D(10, kernel_size=3, activation='relu')
```

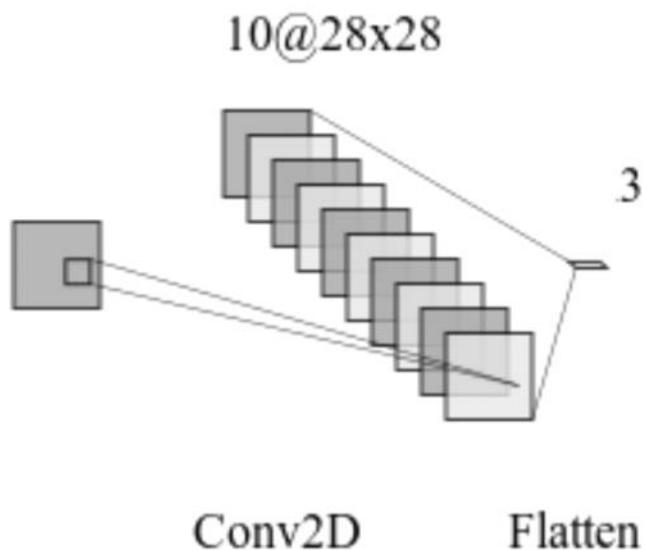
Integrating convolution layers into a network

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten

model = Sequential()
model.add(Conv2D(10, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```



Our CNN



Fitting a CNN

```
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

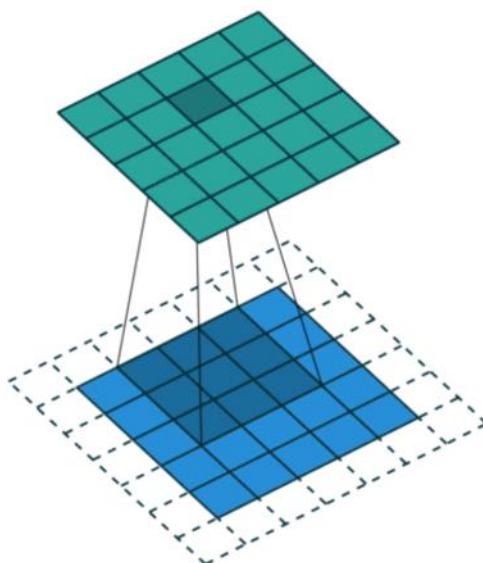
train_data.shape
```

```
(50, 28, 28, 1)
```

```
model.fit(train_data, train_labels, validation_split=0.2,
          epochs=3)

model.evaluate(test_data, test_labels, epochs=3)
```

Convolution with zero padding



Zero padding in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           padding='valid')
```

If we provide the value "valid", no zero padding is added.

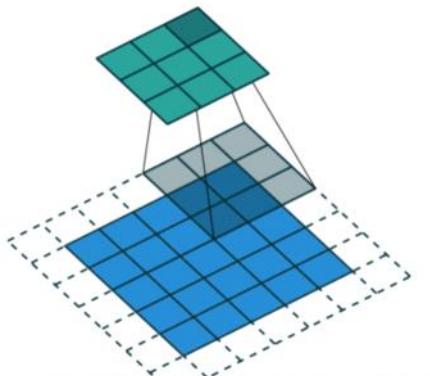
Zero padding in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
            padding='same')
```



this layer, so that the output of the convolution has the same size as the input into the convolution.

Strides



For example, in this animation the kernel is strided by two pixels in each step.

Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
            strides=1)
```

The default is for the stride to be set to 1.

Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           strides=2)
```

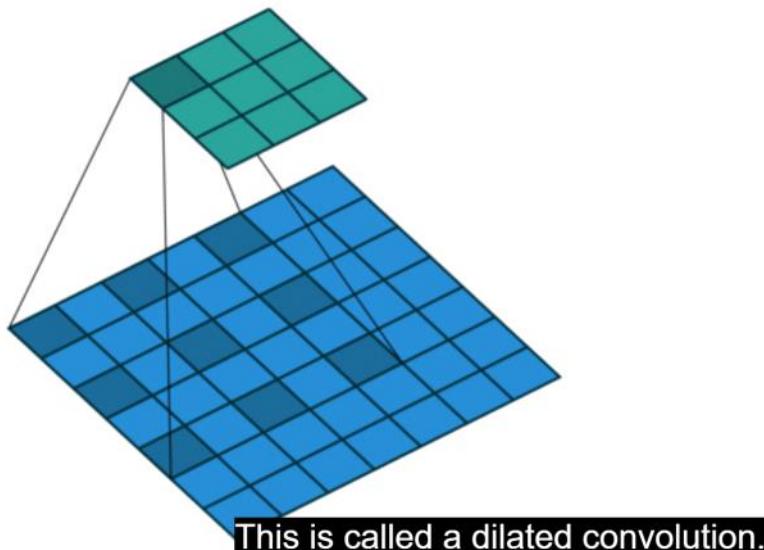
Calculating the size of the output

$$O = ((I - K + 2P)/S) + 1$$

where

- I = size of the input
- K = size of the kernel
- P = size of the zero padding
- S = strides

Dilated convolutions



In this case, the convolution kernel has only 9 parameters, but it

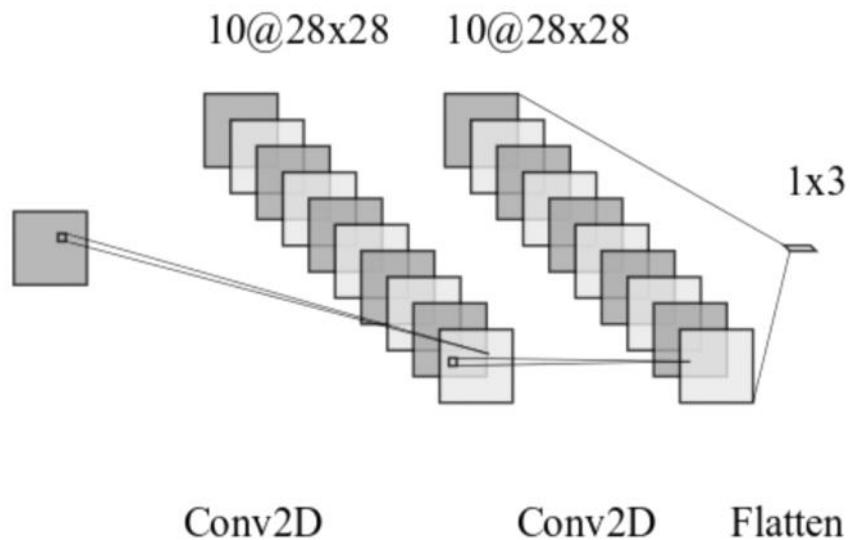
has the same field of view as a kernel that would have the size 5 by 5.

Dilation in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           dilation_rate=2)
```

Going Deeper

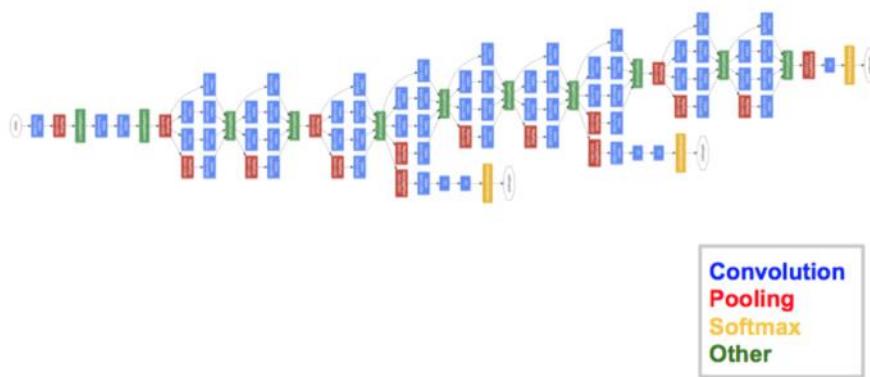
Building a deeper network



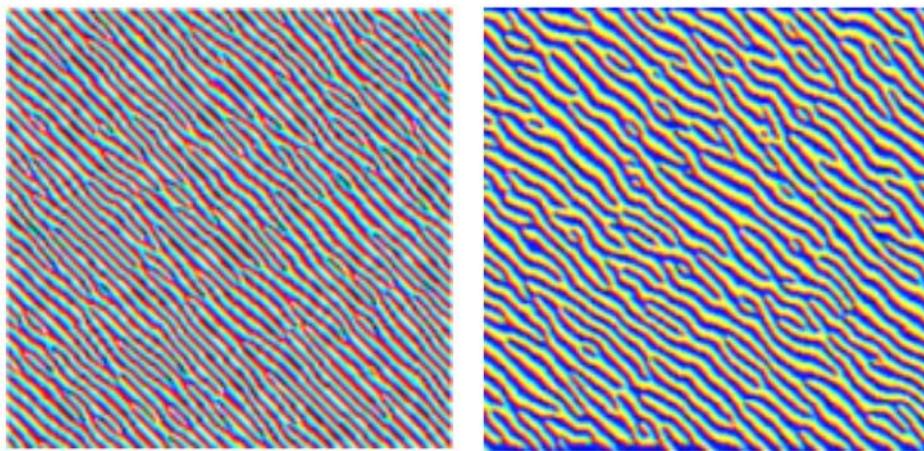
Building a deep network

```
model = Sequential()  
model.add(Conv2D(10, kernel_size=2, activation='relu',  
               input_shape=(img_rows, img_cols, 1),  
               padding='equal'))  
  
# Second convolutional layer  
model.add(Conv2D(10, kernel_size=2, activation='relu'))  
model.add(Flatten())  
model.add(Dense(3, activation='softmax'))
```

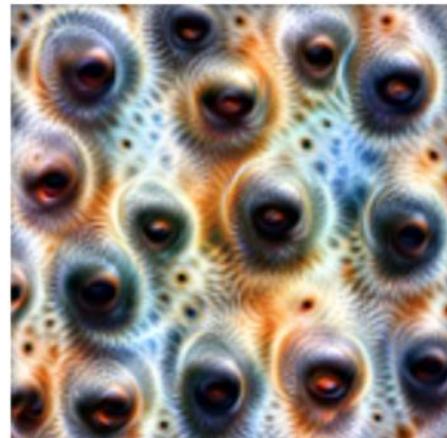
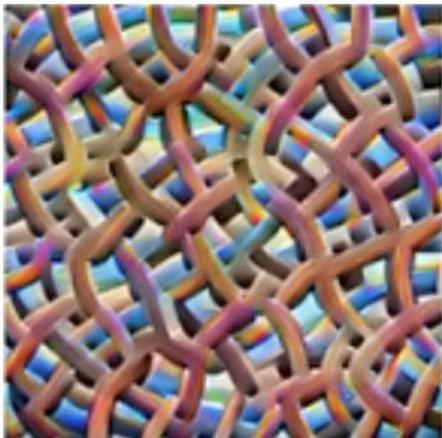
Why do we want deep networks?



Features in early layers



Features in intermediate layers



Features in late layers



How deep?

- Depth comes at a computational cost
- May require more data

Counting parameters

```
model = Sequential()

model.add(Dense(10, activation='relu',
               input_shape=(784, )))

model.add(Dense(10, activation='relu'))

model.add(Dense(3, activation='softmax'))
```

```
# Call the summary method
model.summary()
```

```
-----
Layer (type)          Output Shape       Param #
-----
dense_1 (Dense)      (None, 10)         7850
-----
dense_2 (Dense)      (None, 10)         110
-----
dense_3 (Dense)      (None, 3)          33
-----
Total params: 7,993
Trainable params: 7,993
Non-trainable params: 0
-----
```

Counting parameters

```
model.add(Dense(
    10, activation='relu',
    input_shape=(784, )))
```

parameters = $784 * 10 + 10$
= 7850

in this layer, 10 units + 10 parameters for bias terms in every one of these units

```
model.add(Dense(  
    10, activation='relu'))
```

*parameters = $10 * 10 + 10$
= 110*

In the second layer there is one parameter for

every unit in layer one, times the number of units in this layer + 10 parameters for bias terms.

```
model.add(Dense(  
    3, activation='softmax'))
```

*parameters = $10 * 3 + 3$
= 33*

The last layer has one parameter for every unit in layer 2, times the 3 units in this layer + 3 bias terms.

$$7850 + 110 + 33 = 7993$$

That's why the total number of parameters is 7,993.

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	7850
dense_2 (Dense)	(None, 10)	110
dense_3 (Dense)	(None, 3)	33
Total params:	7,993	
Trainable params:	7,993	
Non-trainable params:	0	

The number of parameters in a CNN

```
model.add(  
    Conv2D(10, kernel_size=3,  
           activation='relu',  
           input_shape=(28, 28, 1),  
           padding='same'))  
parameters = 9 * 10 + 10  
= 100  
  
model.add(  
    Conv2D(10, kernel_size=3,  
           activation='relu',  
           padding='same'))  
parameters = 10 * 9 * 10 + 10  
= 910  
parameters = 0  
  
model.add(Flatten())  
parameters = 7840 * 3 + 3  
= 23523  
model.add(Dense(3))
```

```
model.summary()
```

```
-----  
Layer (type)          Output Shape         Param #  
-----  
conv2d_1 (Conv2D)      (None, 28, 28, 10)      100  
-----  
conv2d_2 (Conv2D)      (None, 28, 28, 10)      910  
-----  
flatten_3 (Flatten)    (None, 7840)            0  
-----  
dense_4 (Dense)        (None, 3)                23523  
-----  
Total params: 24,533  
Trainable params: 24,533  
Non-trainable params: 0
```

Increasing the number of units in each layer

```
model = Sequential()  
  
model.add(Dense(5, activation='relu',  
               input_shape=(784,), padding='same'))  
  
model.add(Dense(15, activation='relu', padding='same'))  
  
model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

```
-----  
Layer (type)          Output Shape         Param #  
=====-----  
dense_1 (Dense)      (None, 5)           3925  
-----  
dense_2 (Dense)      (None, 15)          90  
-----  
dense_3 (Dense)      (None, 3)           48  
-----  
Total params: 4,063  
Trainable params: 4,063  
Non-trainable params: 0  
-----
```

Increasing the number of units in each layer

```
model = Sequential()  
  
model.add(Conv2D(5, kernel_size=3, activation='relu',  
                input_shape=(28, 28, 1),  
                padding="same"))  
  
model.add(Conv2D(15, kernel_size=3, activation='relu',  
                padding="same"))  
  
model.add(Flatten())  
  
model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50
conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
flatten_6 (Flatten)	(None, 11760)	0
dense_9 (Dense)	(None, 3)	35283

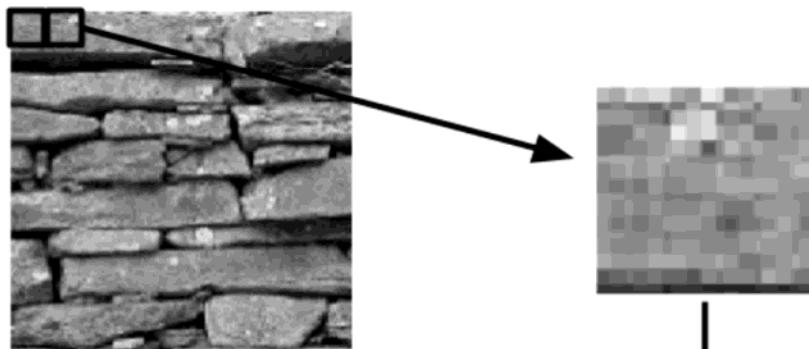
Total params: 36,023
Trainable params: 36,023
Non-trainable params: 0

Pooling operations

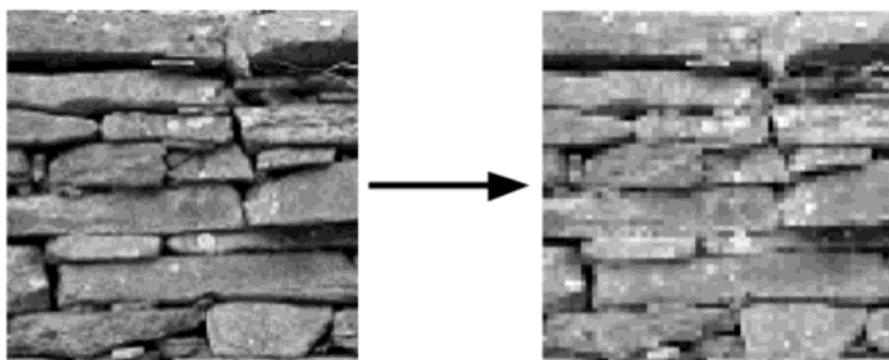
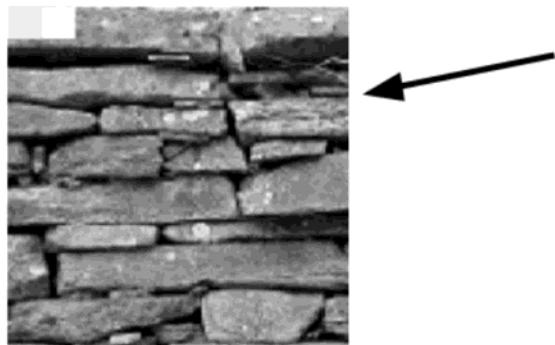
```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50
conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
flatten_6 (Flatten)	(None, 11760)	0
dense_9 (Dense)	(None, 3)	35283

Total params: 36,023
Trainable params: 36,023
Non-trainable params: 0



Max



Implementing max pooling

```
result = np.zeros((im.shape[0]//2, im.shape[1]//2))

result[0, 0] = np.max(im[0:2, 0:2])
result[0, 1] = np.max(im[0:2, 2:4])
result[0, 2] = np.max(im[0:2, 4:6])
```

...

```
result[1, 0] = np.max(im[2:4, 0:2])
result[1, 1] = np.max(im[2:4, 2:4])
```

Implementing max pooling

```
for ii in range(result.shape[0]):
    for jj in range(result.shape[1]):
        result[ii, jj] = np.max(im[ii*2:ii*2+2, jj*2:jj*2+2])
```

Max pooling in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, MaxPool2D

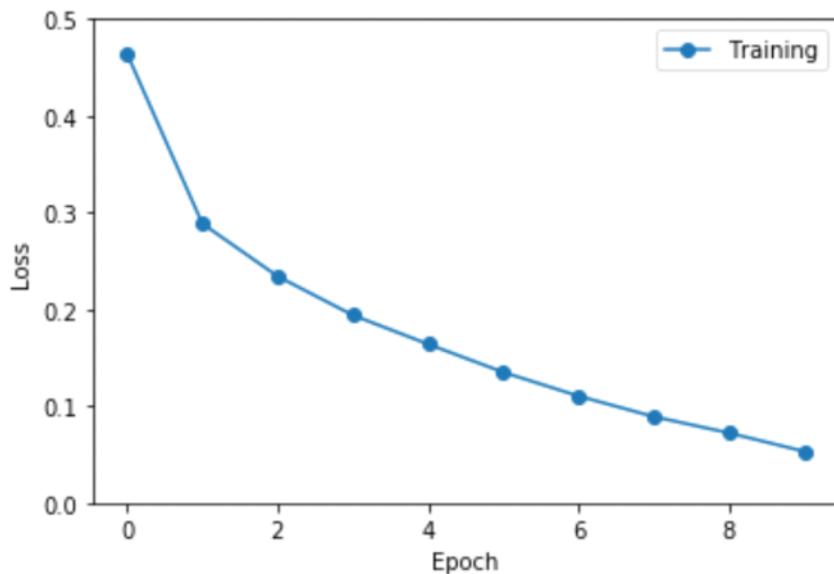
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(MaxPool2D(2))
model.add(Conv2D(15, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(MaxPool2D(2))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

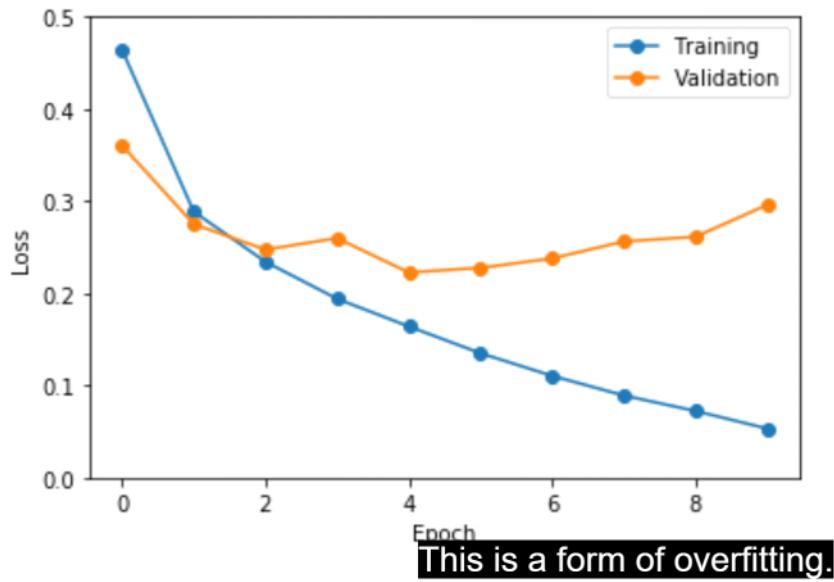
Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 5)	50
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 5)	0
conv2d_2 (Conv2D)	(None, 11, 11, 15)	690
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 15)	0
flatten_1 (Flatten)	(None, 375)	0
dense_1 (Dense)	(None, 3)	1128

Understanding and Improving Deep Convolutional Networks

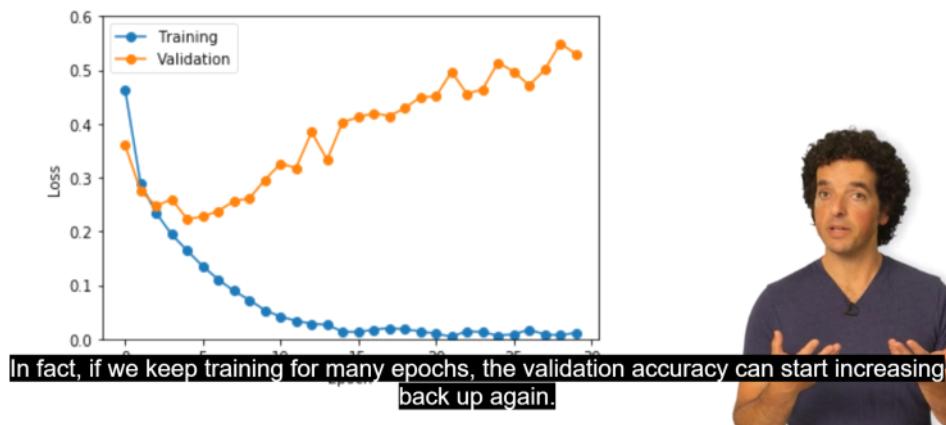
Learning curves: training



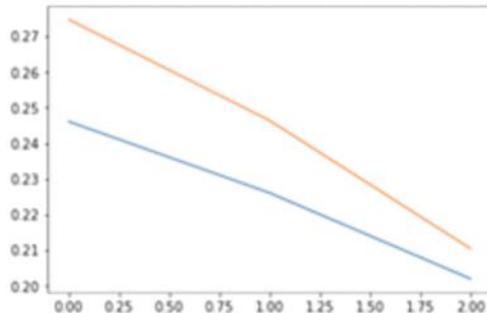
Learning curves: validation



Learning curves: overfitting



```
training = model.fit(train_data, train_labels,  
                      epochs=3, validation_split=0.2)  
  
import matplotlib.pyplot as plt  
plt.plot(training.history['loss'])  
plt.plot(training.history['val_loss'])  
plt.show()
```



Storing the optimal parameters

```
from keras.callbacks import ModelCheckpoint  
  
# This checkpoint object will store the model parameters  
# in the file "weights.hdf5"  
checkpoint = ModelCheckpoint('weights.hdf5', monitor='val_loss',  
                             save_best_only=True)  
  
# Store in a list to be used during training  
callbacks_list = [checkpoint]  
  
# Fit the model on a training set, using the checkpoint as a  
#callback  
model.fit(train_data, train_labels, validation_split=0.2,  
          epochs=3, callbacks=callbacks_list)
```

Loading stored parameters

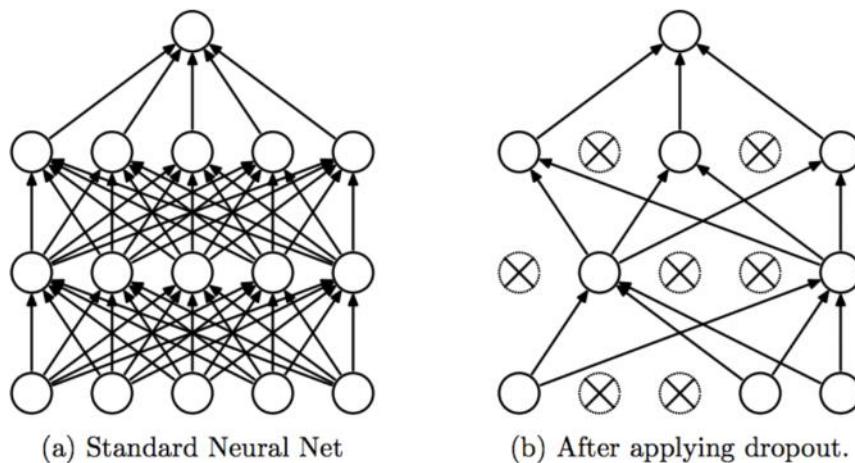
```
model.load_weights('weights.hdf5')  
model.predict_classes(test_data)  
array([2, 2, 1, 2, 0, 1, 0, 1, 2, 0])
```

Dropout

In each learning step:

- Select a subset of the units
- Ignore it in the forward pass
- And in the back-propagation of error

Dropout



Dropout in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, Dropout

model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(Dropout(0.25))

model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Batch normalization

- Rescale the outputs

Batch Normalization in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, BatchNo

model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))

model.add(BatchNormalization())

model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Be careful when using them together!

The disharmony between dropout and batch normalization

This is because while dropout slows down learning, making it more

incremental and careful, batch normalization tends to make learning go faster.

Their effects together may in fact counter each other, and networks sometimes perform

Selecting layers

```
model.layers
```

Once a model is constructed and compiled, it will store its layers in an attribute called "layers".

```
[<keras.layers.convolutional.Conv2D at 0x109f10c18>,
 <keras.layers.convolutional.Conv2D at 0x109ec5ba8>,
 <keras.layers.core.Flatten at 0x1221ffcc0>,
 <keras.layers.core.Dense at 0x1221ffef0>]
```

Getting model weights

```
conv1 = model.layers[0]
weights1 = conv1.get_weights()
len(weights1)
```

```
2
```

```
kernels1 = weights1[0]
kernels1.shape
```

```
(3, 3, 1, 5)
```

The first 2 dimensions denote the kernel size.

The third dimension denotes the number of channels in the kernels.

The last dimension denotes the number of kernels in this layer: 5.

```
kernel1_1 = kernels1[:, :,  
                      0, 0]
```

```
kernel1_1.shape
```

(3, 3)

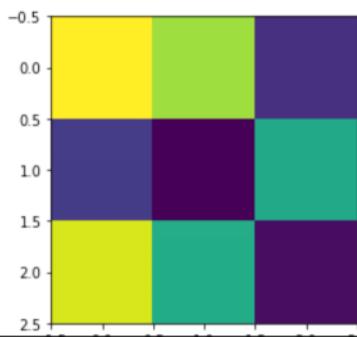
To pull out the first kernel in this layer, we would use the index 0 into the last dimension

Because there is only one channel, we can also index on the channel dimension, to collapse that dimension.

This would return the 3 by 3 array containing this convolutional kernel.

Visualizing the kernel

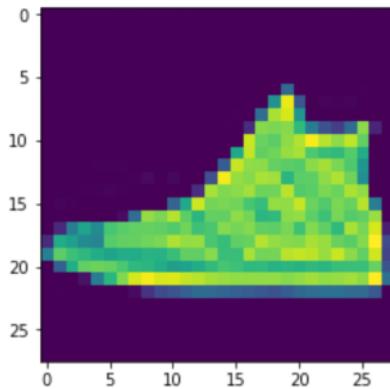
```
plt.imshow(kernel1_1)
```



We can then visualize this kernel directly, but understanding what kinds of features this kernel is responding to may be hard just from direct observation

Visualizing the kernel responses

```
test_image = test_data[3, :, :, 0]
plt.imshow(test_image)
```

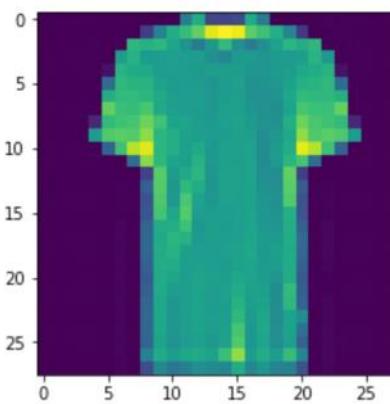


Visualizing the kernel responses

```
filtered_image = convolution(test_image, kernel1_1)
plt.imshow(filtered_image)
```

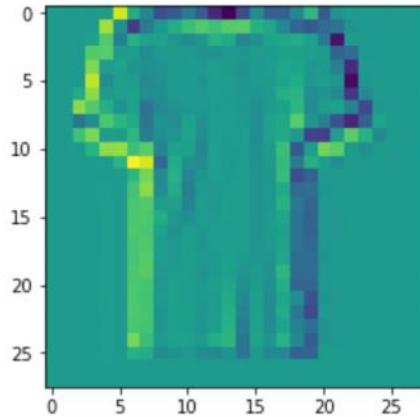
Visualizing the kernel responses

```
test_image = test_data[4, :, :, 1]
plt.imshow(test_image)
```



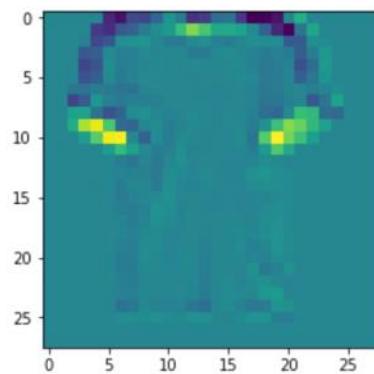
Visualizing the kernel responses

```
filtered_image = convolution(test_image, kernel1_1)  
plt.imshow(filtered_img)
```



Visualizing the kernel responses

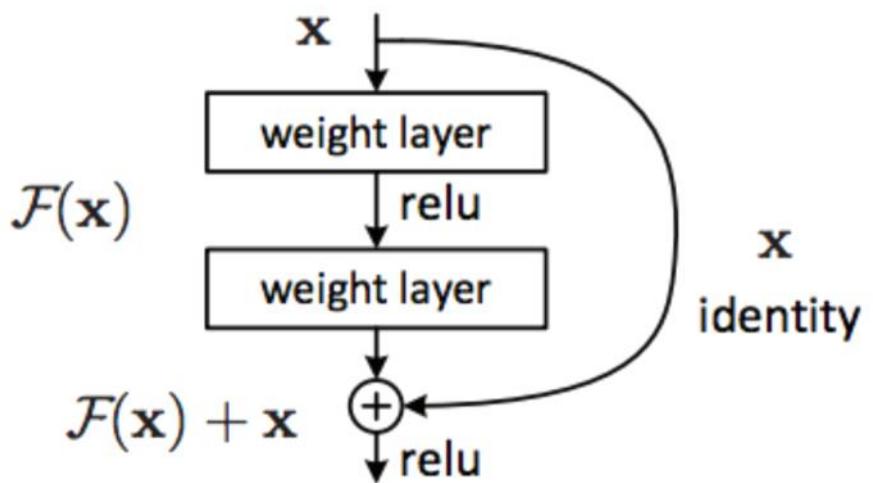
```
kernel1_2 = kernels[:, :, 0, 1]  
filtered_image = convolution(test_image, kernel1_2)  
plt.imshow(filtered_img)
```



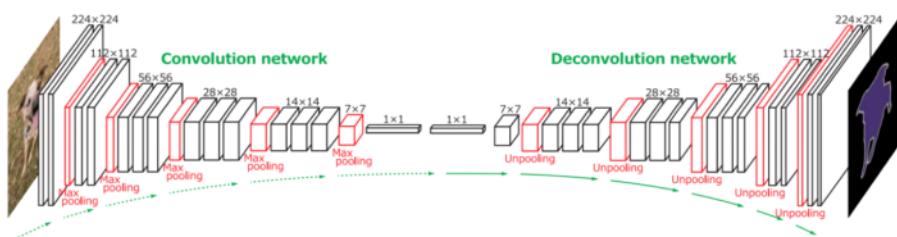
What next?

- Even deeper networks
- Residual networks
- Transfer learning
- Fully convolutional networks

Residual networks



Fully convolutional networks



Generative adversarial networks

